

COMP0086 : Summative Assessment

University College London

13 November 2024

Table of Contents

1	Question 1 - Models for Binary Vectors	2
1.1	Question 1.(a)	2
1.2	Question 1.(b)	2
1.3	Question 1.(c)	3
1.4	Question 1.(d)	5
1.5	Question 1.(e)	9
2	Question 2 - Model selection	10
2.1	Question 2.(a)	10
2.2	Question 2.(b)	11
2.3	Question 2.(c)	12
2.4	Question 2. Answer	13
3	Question 3 - EM for Binary Data	18
3.1	Question 3.(a)	18
3.2	Question 3.(b)	20
3.3	Question 3.(c)	21
3.4	Question 3.(d)	23
3.5	Question 3.(e)	31
4	Question 5 - Decrypting Messages with MCMC	34
4.1	Question 5.(a)	34
4.2	Question 5.(b)	37
4.3	Question 5.(c)	37
4.4	Question 5.(d)	38
4.5	Question 5.(e)	49
4.6	Question 5.(f)	49
5	Question 7 - Optimization	51
5.1	Question 7.(a)	51
5.2	Question 7.(b)	52
5.2.1	Question 7.(b).(i)	52
5.2.2	Question 7.(b).(ii)	52
6	Question 8 - [BONUS] Eigenvalues as solutions of an optimization problem	53
6.1	Question 8.(a)	53
6.2	Question 8.(b)	54
6.3	Question 8.(c)	56

1 | Question 1 - Models for Binary Vectors

1.1 Question 1.(a)

We dispose of a data set of N binary images. Each of these images has a finite number of D discrete pixels, each taking a value 0 or 1, but not in between. The Multivariate Gaussian distribution is defined on the domain $x \in \mathbb{R}^k$, and its data type is continuous and unbounded. Compared to our data which is discrete and bounded ($x_d^{(n)} \in \{0, 1\}$), then the Multivariate Gaussian distribution is not suited for the data set of images.

1.2 Question 1.(b)

We assume the images were modelled as independently and identically distributed (**iid**) samples from a D -dimensional multivariate Bernoulli distribution with parameter vector $\mathbf{p} = (p_1, \dots, p_D)$:

$$P(\mathbf{x}^{(n)}|\mathbf{p}) = \prod_{d=1}^D p_d^{x_d^{(n)}} (1 - p_d)^{1-x_d^{(n)}} \quad n \in 1, \dots, N \quad (1.1)$$

where \mathbf{x} , \mathbf{p} are D -dimensional vectors.

The equation for the likelihood of \mathbf{p} is calculated as:

$$P(\mathbf{x}|\mathbf{p}) = \prod_{n=1}^N P(\mathbf{x}^{(n)}|\mathbf{p}) = \prod_{n=1}^N \prod_{d=1}^D p_d^{x_d^{(n)}} (1 - p_d)^{1-x_d^{(n)}} \quad (1.2)$$

To obtain the log-likelihood, we take the log of the expression, yielding:

$$\begin{aligned} \mathcal{L} &= \log \prod_{n=1}^N \prod_{d=1}^D p_d^{x_d^{(n)}} (1 - p_d)^{1-x_d^{(n)}} \\ \mathcal{L} &= \sum_{n=1}^N \sum_{d=1}^D p_d^{x_d^{(n)}} (1 - p_d)^{1-x_d^{(n)}} \\ \mathcal{L} &= \sum_{n=1}^N \sum_{d=1}^D x_d^{(n)} \log p_d + (1 - x_d^{(n)}) \log (1 - p_d) \end{aligned} \quad (1.3)$$

To evaluate the Maximum Likelihood, we need to find the maximum of that function:

$$\frac{\partial \mathcal{L}(p)}{\partial p_d} = 0$$

Equivalent to:

$$\sum_{n=1}^N \left(\frac{x_d^{(n)}}{p_d} - \frac{(1 - x_d^{(n)})}{1 - p_d} \right) = 0$$

$$\sum_{n=1}^N \left(\frac{x_d^{(n)}(1-p_d) - p_d(1-x^{(n)})}{p_d(1-p_d)} \right) = 0$$

Which is achieved when the numerator is equal to 0:

$$\sum_{n=1}^N x_d^{(n)}(1-p_d) - p_d(1-x^{(n)}) = 0$$

$$\sum_{n=1}^N (x_d^{(n)} - p_d) = 0$$

$$\sum_{n=1}^N x_d^{(n)} = \sum_{n=1}^N p_d$$

And since p_d does not depend on N , we obtain the Maximum Likelihood (ML) estimation of p_d , denoted as \hat{p}_d^{ML} :

$$\begin{aligned} \sum_{n=1}^N x_d^{(n)} &= Np_d \\ \hat{p}_d^{ML} &= \frac{1}{N} \sum_{n=1}^N x^{(n)} \end{aligned} \tag{1.4}$$

Generalizing for all values of d since the pixels are assumed to be independent, we get:

$$\hat{\mathbf{p}}^{\text{ML}} = \frac{1}{N} \sum_{n=1}^N \mathbf{x}^{(n)} \tag{1.5}$$

1.3 Question 1.(c)

We assume independent Beta priors on the parameters p_d :

$$P(p_d) = \frac{1}{B(\alpha, \beta)} p_d^{\alpha-1} (1-p_d)^{\beta-1} \tag{1.6}$$

And:

$$P(\mathbf{p}) = \prod_d P(p_d) \tag{1.7}$$

To find the MAP estimator of \mathbf{p} , we firstly use Bayes Rule:

$$P(\mathbf{p}|\mathbf{x}) = \frac{P(\mathbf{x}|\mathbf{p})P(\mathbf{p})}{P(\mathbf{x})} \tag{1.8}$$

Combining the expressions, we get:

$$P(\mathbf{p}|\mathbf{x}) = \frac{\prod_{n=1}^N \prod_{d=1}^D p_d^{x_d^{(n)}} (1-p_d)^{1-x_d^{(n)}} \prod_{d=1}^D P(p_d)}{P(\mathbf{x})} \tag{1.9}$$

The expression is hard to compute, but, since $P(\mathbf{x})$ does not depend on d , we can state the following:

$$P(\mathbf{p}|\mathbf{x}) \propto p_d^{x_d^{(n)}} (1 - p_d)^{1-x_d^{(n)}} \quad (1.10)$$

Taking the log of the expression:

$$\begin{aligned} \log P(\mathbf{p}|\mathbf{x}) &= \sum_{n=1}^N \sum_{d=1}^D \left(x_d^{(n)} \log p_d + (1 - x_d^{(n)}) \log (1 - p_d) \right) \\ &+ \sum_{d=1}^D (\log(p_d)(\alpha - 1) + (\beta - 1) \log (1 - p_d)) - D \log B(\alpha, \beta) \end{aligned}$$

We need to find the maximum of this expression. Differentiating with respect to p_d , we get:

$$\frac{\partial \log P(\mathbf{p}|\mathbf{x})}{\partial p_d} = \sum_{n=1}^N \left(x_d^{(n)} \frac{1}{p_d} + (1 - x_d^{(n)}) \frac{-1}{1 - p_d} \right) + \frac{(\alpha - 1)}{p_d} - \frac{\beta - 1}{1 - p_d} \quad (1.11)$$

Because the Beta function term is independent of d .

$$\begin{aligned} \frac{\partial \log P(\mathbf{p}|\mathbf{x})}{\partial p_d} &= \sum_{n=1}^N \left(\frac{x_d^{(n)}(1 - p_d) - (1 - x_d^{(n)})p_d}{p_d(1 - p_d)} \right) + \frac{(\alpha - 1)(1 - p_d) - (\beta - 1)p_d}{p_d(1 - p_d)} \\ &= \sum_{n=1}^N \left(\frac{x_d^{(n)} - p_d - x_d^{(n)}p_d + x_d^{(n)}p_d}{p_d(1 - p_d)} \right) + \frac{\alpha - 1 - \alpha p_d + p_d - \beta p_d + p_d}{p_d(1 - p_d)} \\ \frac{\partial \log P(\mathbf{p}|\mathbf{x})}{\partial p_d} &= \frac{\sum_{n=1}^N (x_d^{(n)} - p_d) + \alpha - 1 + p_d(2 - \alpha - \beta)}{p_d(1 - p_d)} \end{aligned} \quad (1.12)$$

Making this expression equal to 0:

$$\frac{\partial \log P(\mathbf{p}|\mathbf{x})}{\partial p_d} = 0 \quad (1.13)$$

Equivalent to:

$$\sum_{n=1}^N (x_d^{(n)} - p_d) + \alpha - 1 + p_d(2 - \alpha - \beta) = 0 \quad (1.14)$$

$$\sum_{n=1}^N (x_d^{(n)}) - Np_d + \alpha - 1 + p_d(2 - \alpha - \beta) = 0$$

$$p_d(N + \alpha + \beta - 2) = \sum_{n=1}^N (x_d^{(n)}) + \alpha - 1$$

$$\hat{p}_d^{MAP} = \frac{\sum_{n=1}^N (x_d^{(n)}) + \alpha - 1}{N + \alpha + \beta - 2} \quad (1.15)$$

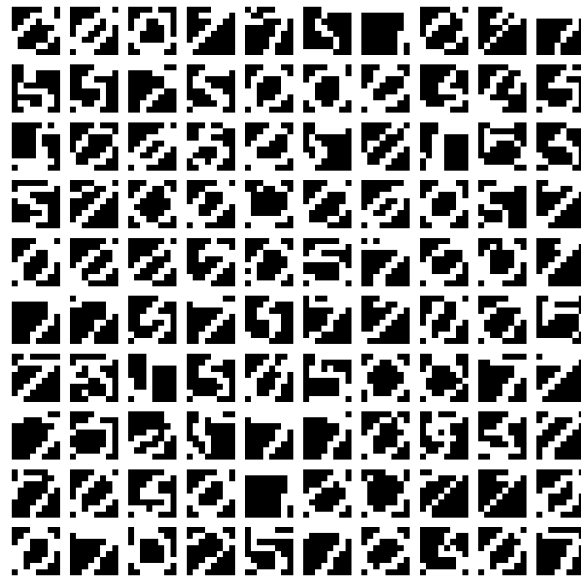


Figure 1.1: Original data from [binarydigits.txt](#), used for training

And thus, generalizing for \mathbf{p} , we obtain the Maximum A Posteriori (MAP) expression of \mathbf{p} :

$$\hat{\mathbf{p}}^{MAP} = \frac{\sum_{n=1}^N (x^{(n)}) + \alpha - 1}{N + \alpha + \beta - 2} \quad (1.16)$$

1.4 Question 1.(d)

The original data contains $N = 100$ image, each made of $D = 64$ pixels and stored in an $N \times D$ matrix. Rearranging the pixels, Figure 1.1 shows the original data from the [binarydigits.txt](#) file, and displaying them as an 8×8 image, which was displayed by readapting the code from [bindigit.py](#).

After executing the code to learn the ML parameters of the multivariate Bernoulli from the dataset, the obtained ML parameters are shown in Figure 1.2(a). Those parameters are displayed as an 8×8 image. The [code](#) for the execution can be found below.

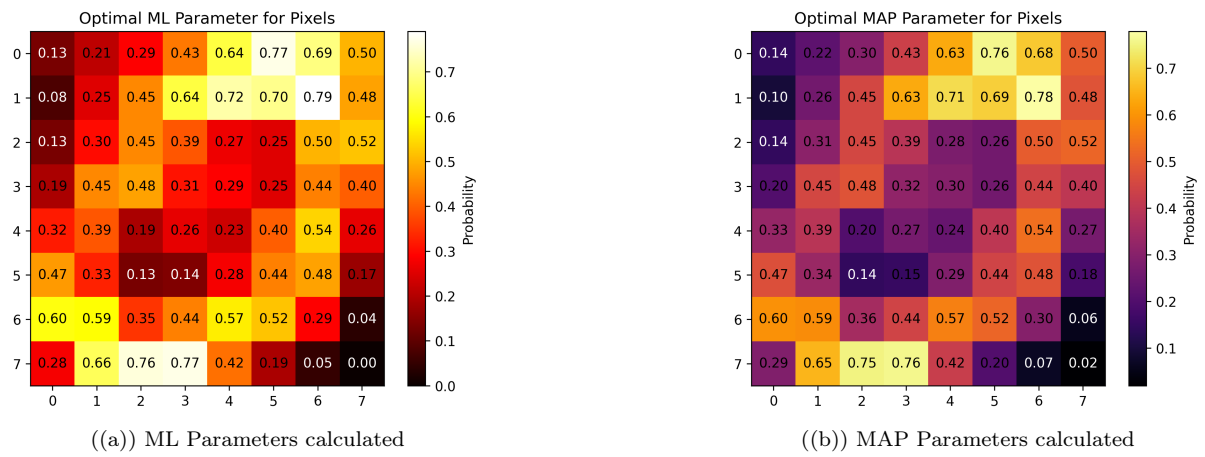


Figure 1.2: ML and MAP Parameters

Listing 1: Code - Q1.(d)

```
1 import numpy as np
2 from matplotlib import pyplot as plt
3
4 # Loading the training data
5 def load_binarydigits(filename='binarydigits.txt'):
6     Y = np.loadtxt('binarydigits.txt')
7     return Y
8
9 # Displaying the original training data
10 def save_data_binary(Y: np.ndarray, save=False):
11     N, _ = Y.shape
12     plt.figure(figsize=(5, 5))
13     for n in range(N):
14         plt.subplot(10, 10, n+1)
15         plt.imshow(np.reshape(Y[n, :], (8,8)),
16                     interpolation="None",
17                     cmap='gray')
18         plt.axis('off')
19     if save:
20         plt.savefig('data_binary.png', format="png", dpi=300,
21                     bbox_inches="tight")
22     plt.show()
23
24 # ML Parameter Learning
25 def ML_learning(Y: np.ndarray, save_ML = False):
26     N, D = Y.shape
27     p_ML = np.zeros((D, 1), dtype=np.float64)
28
29     for d in range(D):
30         p_ML[d] = (1/N)*np.sum(Y[:, d])
31
32     p_ML_image = np.reshape(p_ML, (8,8))
33     plt.figure()
34     plt.imshow(p_ML_image, cmap="hot", interpolation='nearest')
35     plt.colorbar(label='Probability')
36     plt.title('Optimal ML Parameter for Pixels')
37     for i in range(8):
38         for j in range(8):
39             if p_ML_image[i,j] > 0.15:
40                 plt.text(j, i, f"{p_ML_image[i, j]:.2f}", ha='center',
41                         va='center', color="black")
42             else:
43                 plt.text(j, i, f"{p_ML_image[i, j]:.2f}", ha='center',
44                         va='center', color="white")
45     if save_ML:
46         plt.savefig('ML_parameter.png', format="png", dpi=300, bbox_inches="tight")
47     plt.show()
```



```
48         return p_ML_image, p_ML
49
50
51 def main():
52     Y = load_binarydigits()
53     save_data_binary(Y)
54     p_ML_image, p_ML = ML_learning(Y)
55
56     # Executing the main function
57     main()
```

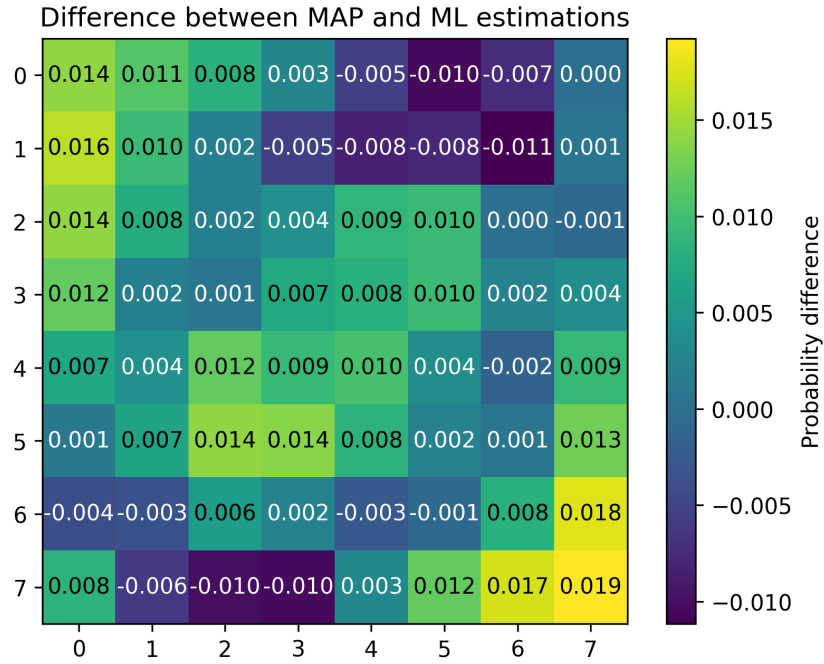


Figure 1.3: Subtracting ML parameters to MAP parameters: $\hat{\mathbf{p}}^{MAP} - \hat{\mathbf{p}}^{ML}$

1.5 Question 1.(e)

Modifying the code built for the previous question, we obtain the MAP parameters displayed in Figure 1.2(b), which appear different if we look closely to the displayed numbers on each tile and compare them with that of the ML parameters. For the difference to be more visible, Figure 1.3 shows the difference between MAP and ML parameters, calculated as $\hat{\mathbf{p}}^{MAP} - \hat{\mathbf{p}}^{ML}$ ¹. It provides a better visual of the difference between these two estimations.

The new code is displayed below.

The new learned parameters are better than the ML estimate due to the limited amount of data. If we had more data, the MAP estimate would yield the same result as the ML estimate. Having the MAP estimate permits the model to not assume some pixels are 1's and others are 0's by defaults, because it takes into account what it has seen on the prior.

¹Why the difference, and not the quotient: some entries are 0, the computer will display an error in such case.

2 | Question 2 - Model selection

We would like to find the expressions needed to calculate the relative probabilities of three different models regarding our binary data images.

2.1 Question 2.(a)

For this model (named *Model 1*), we assume all D components are generated from a Bernoulli distribution with identical $p_d = 0.5$. Thus, the probability of data x given $p_d = 0.5$ is:

$$P_1(x^n | \mathbf{p} = [p_d, \dots, p_d]^T, p_d = 0.5) = \prod_{i=1}^D p_d^{x_i^{(n)}} (1 - p_d)^{1-x_i^{(n)}} \quad (2.1)$$

$$= \prod_{d=i}^D 0.5^{x_i^{(n)}} (1 - 0.5)^{1-x_i^{(n)}}$$

$$= \prod_{i=1}^D 0.5^{x_i^{(n)} + 1 - x_i^{(n)}}$$

$$= \prod_{i=1}^D 0.5$$

$$P_1(x^{(n)} | \mathbf{p} = [p_d, \dots, p_d]^T, p_d = 0.5) = 0.5^D \quad (2.2)$$

Thus, for the entire dataset, since all images are considered to be independent and identically sampled (iid):

$$P_1(\mathbf{x} | \mathbf{p} = [p_d, \dots, p_d]^T, p_d = 0.5) = \prod_{n=1}^N P_1(x^{(n)} | \mathbf{p} = [p_d, \dots, p_d]^T, p_d = 0.5) \quad (2.3)$$

$$= \prod_{n=1}^N 0.5^D$$

$$P_1(\mathbf{x} | \mathbf{p} = [p_d, \dots, p_d]^T, p_d = 0.5) = 0.5^{N \times D} \quad (2.4)$$

For each model, the likelihood of the data \mathbf{x} given the model i is given by:

$$P(\mathbf{x} | \text{Model } i) = \int_0^1 P_i(\mathbf{x} | \mathbf{p}) P(\mathbf{p}) d\mathbf{p} \quad (2.5)$$

Where $P(\mathbf{p})$ is known to be uniform for all unknown probabilities. In the case of Model 1, we know the prior probability and it is equal to 1 (no probability of having something else than $p_d = 0.5$). Therefore:

$$P(\mathbf{x} | \text{Model 1}) = \int_0^1 P_1(\mathbf{x} | \mathbf{p}) d\mathbf{p} \quad (2.6)$$

$$= \int_0^1 0.5^{N \times D} d\mathbf{p}$$

$$= 0.5^{N \times D} \int_0^1 d\mathbf{p}$$

$$P(\mathbf{x}|\text{Model 1}) = 0.5^{N \times D} \quad (2.7)$$

2.2 Question 2.(b)

For this model (named *Model 2*), we assume all D components are generated from Bernoulli distributions with unknown but identical p_d . The probability $P_2(x^{(n)}|\mathbf{p} = [p_d, \dots, p_d]^T)$ is:

$$P_2(x^{(n)}|\mathbf{p}) = \prod_{i=1}^D p_d^{x_i^{(n)}} (1 - p_d)^{1-x_i^{(n)}} \quad (2.8)$$

Leading to:

$$P_2(x^{(n)}|\mathbf{p}) = p_d^{\sum_{i=1}^D x_i^{(n)}} (1 - p_d)^{\sum_{i=1}^D (1-x_i^{(n)})} \quad (2.9)$$

N.B. notice we changed the indice of the product, to avoid confusion.
For the entire image dataset, we obtain:

$$P_2(\mathbf{x}|\mathbf{p}) = \prod_{n=1}^N P_2(x^{(n)}|\mathbf{p}) \quad (2.10)$$

Which we calculated before:

$$\begin{aligned} &= \prod_{n=1}^N p_d^{\sum_{i=1}^D x_i^{(n)}} (1 - p_d)^{\sum_{i=1}^D (1-x_i^{(n)})} \\ &= p_d^{\sum_{n=1}^N \sum_{i=1}^D x_i^{(n)}} (1 - p_d)^{\sum_{n=1}^N \sum_{i=1}^D (1-x_i^{(n)})} \end{aligned}$$

Thus, as for Model 1, to obtain the likelihood of Model 2:

$$P(\mathbf{x}|\text{Model 2}) = \int_0^1 P_2(\mathbf{x}|\mathbf{p}) P(\mathbf{p}) d\mathbf{p} \quad (2.11)$$

And since we assumed a uniform prior because p_d is unknown, therefore it becomes:

$$\begin{aligned} P(\mathbf{x}|\text{Model 2}) &= \int_0^1 P_2(\mathbf{x}|\mathbf{p}) dp_d \\ &= \int_0^1 p_d^{\sum_{n=1}^N \sum_{i=1}^D x_i^{(n)}} (1 - p_d)^{\sum_{n=1}^N \sum_{i=1}^D (1-x_i^{(n)})} dp_d \end{aligned} \quad (2.12)$$

And since integrating a Binomial function (parameters over the interval $[0, 1]$ is the same as evaluating the Beta function:

$$B(\alpha, \beta) = \int_0^1 p^{\alpha-1} (1-p)^{\beta-1} dp$$

Then, by identification:

$$P(\mathbf{x}|\text{Model 2}) = B\left(1 + \sum_{n=1}^N \sum_{i=1}^D x_i^{(n)}, 1 + \sum_{n=1}^N \sum_{i=1}^D (1-x_i^{(n)})\right)$$

Moreover, the Beta function, if its parameters are integers, can be evaluated as:

$$\text{Beta}(i, j) = \frac{(i-1)!(j-1)!}{(i+j-1)!}$$

Since $x_i^{(n)}$'s only take 0's or 1's as values, we can deduce their sums will be integers and therefore:

$$P(\mathbf{x}|\text{Model 2}) = \frac{\left(\sum_{n=1}^N \sum_{i=1}^D x_i^{(n)}\right)! \left(\sum_{n=1}^N \sum_{i=1}^D (1 - x_i^{(n)})\right)!}{\left(\sum_{n=1}^N \sum_{i=1}^D x_i^{(n)} + \sum_{n=1}^N \sum_{i=1}^D (1 - x_i^{(n)}) + 1\right)!} \quad (2.13)$$

Finally, we can rapidly develop the second double sum to obtain:

$$P(\mathbf{x}|\text{Model 2}) = \frac{\left(\sum_{n=1}^N \sum_{i=1}^D x_i^{(n)}\right)! \left(N \times D + \sum_{n=1}^N \sum_{i=1}^D (-x_i^{(n)})\right)!}{\left(\sum_{n=1}^N \sum_{i=1}^D x_i^{(n)} + N \times D + \sum_{n=1}^N \sum_{i=1}^D (-x_i^{(n)}) + 1\right)!}$$

Leading to the final expression:

$$P(\mathbf{x}|\text{Model 2}) = \frac{\left(\sum_{n=1}^N \sum_{i=1}^D x_i^{(n)}\right)! \left(N \times D - \sum_{n=1}^N \sum_{i=1}^D (x_i^{(n)})\right)!}{(N \times D + 1)!} \quad (2.14)$$

A note on implementation After implementing and testing the model's probability, the difficulty of implementing large factorial functions (displaying errors due to too intensive computation) led to finally implement them using the log of Beta function formula. This have been done using *betaln()* function from *scipy.special* library.

2.3 Question 2.(c)

We now consider a model (named *Model 3*) with each component being Bernoulli distributed, with a separate and unknown p_d . We therefore have:

$$P_3(x^{(n)}|\mathbf{p}) = \prod_{d=1}^D p_d^{x_d^{(n)}} (1 - p_d)^{(1-x_d^{(n)})} \quad (2.15)$$

$$P_3(\mathbf{x}|\mathbf{p}) = \prod_{n=1}^N \prod_{d=1}^D p_d^{x_d^{(n)}} (1 - p_d)^{(1-x_d^{(n)})} \quad (2.16)$$

For Model 3's likelihood, it is obtained similarly to that of Model 2:

$$P(\mathbf{x}|\text{Model 3}) = \int_0^1 P_3(\mathbf{x}|\mathbf{p}) P(\mathbf{p}) d\mathbf{p} \quad (2.17)$$

But since each component has a separate unknown p_d , still with uniform prior, we get:

$$P(\mathbf{x}|\text{Model 3}) = \int_0^1 \dots \int_0^1 P_3(\mathbf{x}|\mathbf{p}) dp_1 \dots dp_D \quad (2.18)$$

Equivalent to:

$$P(\mathbf{x}|\text{Model 3}) = \int_0^1 \dots \int_0^1 \prod_{n=1}^N \prod_{d=1}^D p_d^{x_d^{(n)}} (1 - p_d)^{(1-x_d^{(n)})} dp_1 \dots dp_D$$

Which can be separated into (since all p_d are independent and separate):

$$\begin{aligned} P(\mathbf{x}|\text{Model 3}) &= \prod_{d=1}^D \left(\int_0^1 \prod_{n=1}^N p_d^{x_d^{(n)}} (1 - p_d)^{(1-x_d^{(n)})} dp_d \right) \\ &= \prod_{d=1}^D \left(\int_0^1 p_d^{\sum_{n=1}^N x_d^{(n)}} (1 - p_d)^{\sum_{n=1}^N (1-x_d^{(n)})} dp_d \right) \end{aligned}$$

Leading to:

$$\begin{aligned} P(\mathbf{x}|\text{Model 3}) &= \prod_{d=1}^D B \left(1 + \sum_{n=1}^N x_d^{(n)}, 1 + \sum_{n=1}^N (1 - x_d^{(n)}) \right) \\ P(\mathbf{x}|\text{Model 3}) &= \prod_{d=1}^D \frac{(\sum_{n=1}^N x_d^{(n)})! (\sum_{n=1}^N (1 - x_d^{(n)}))!}{(\sum_{n=1}^N x_d^{(n)} + \sum_{n=1}^N (1 - x_d^{(n)}) + 1)!} \end{aligned} \quad (2.19)$$

And, after manipulations, leads us to the final expression:

$$\begin{aligned} P(\mathbf{x}|\text{Model 3}) &= \prod_{d=1}^D \frac{(\sum_{n=1}^N x_d^{(n)})! (\sum_{n=1}^N (1 - x_d^{(n)}))!}{(N + 1)!} \\ P(\mathbf{x}|\text{Model 3}) &= \frac{1}{D(N + 1)!} \times \prod_{d=1}^D \left(\sum_{n=1}^N x_d^{(n)} \right)! \left(\sum_{n=1}^N (1 - x_d^{(n)}) \right)! \end{aligned} \quad (2.20)$$

A note on implementation Again, similarly to Model 2's implementation, the difficulty of implementing large factorial functions for Model 3 led to finally implement them using the log of Beta function formula. This have been done using `betaln()` function from `scipy.special` library (once again).

2.4 Question 2. Answer

We assume all models 1, 2 and 3 are equally likely *a priori*, and their prior distributions to be uniform for any unknown probabilities. We would like to find the posterior probabilities of each of the three models.

From Bayes Rule, $i \in 1, 2, 3$:

$$P(\text{Model } i|\mathbf{x}) = \frac{P(\mathbf{x}|\text{Model } i)P(\text{Model } i)}{P(\mathbf{x})} \quad (2.21)$$

Since all three models are equally likely *a priori*:

$$P(\text{Model 1}) = P(\text{Model 2}) = P(\text{Model 3}) = \frac{1}{3}$$

And using:

$$P(\mathbf{x}) = P(\mathbf{x}, \text{Model 1}) + P(\mathbf{x}, \text{Model 2}) + P(\mathbf{x}, \text{Model 3})$$

$$P(\mathbf{x}) = P(\mathbf{x}|\text{Model 1})P(\text{Model 1}) + P(\mathbf{x}|\text{Model 2})P(\text{Model 2}) + P(\mathbf{x}|\text{Model 3})P(\text{Model 3})$$

Leading to:

$$P(\mathbf{x}) = \frac{1}{3} (P(\mathbf{x}|\text{Model 1}) + P(\mathbf{x}|\text{Model 2}) + P(\mathbf{x}|\text{Model 3})) \quad (2.22)$$

We therefore get for Model 1:

$$P(\text{Model 1}|\mathbf{x}) = \frac{P(\mathbf{x}|\text{Model 1})P(\text{Model 1})}{P(\mathbf{x})}$$

Simplifying to (due to equally likely models):

$$P(\text{Model 1}|\mathbf{x}) = \frac{P(\mathbf{x}|\text{Model 1}) \times \frac{1}{3}}{\frac{1}{3} (P(\mathbf{x}|\text{Model 1}) + P(\mathbf{x}|\text{Model 2}) + P(\mathbf{x}|\text{Model 3}))}$$

Permits us to obtain the **MAP equations for Model 1**:

$$P(\text{Model 1}|\mathbf{x}) = \frac{P(\mathbf{x}|\text{Model 1})}{P(\mathbf{x}|\text{Model 1}) + P(\mathbf{x}|\text{Model 2}) + P(\mathbf{x}|\text{Model 3})} \quad (2.23)$$

Similarly for models 2 and 3:

$$P(\text{Model 2}|\mathbf{x}) = \frac{P(\mathbf{x}|\text{Model 2})}{P(\mathbf{x}|\text{Model 1}) + P(\mathbf{x}|\text{Model 2}) + P(\mathbf{x}|\text{Model 3})} \quad (2.24)$$

And

$$P(\text{Model 3}|\mathbf{x}) = \frac{P(\mathbf{x}|\text{Model 3})}{P(\mathbf{x}|\text{Model 1}) + P(\mathbf{x}|\text{Model 2}) + P(\mathbf{x}|\text{Model 3})} \quad (2.25)$$

Log of these probabilities To be easily implemented, the function *betaln()* was the simplest option. However, it is not explicit in the formulas above how they relate. Therefore, for each model i , the probability can be expressed:

$$\begin{aligned} \log P(\text{Model } i|\mathbf{x}) &= \log P(\mathbf{x}|\text{Model } i) - \log (P(\mathbf{x}|\text{Model 1}) + P(\mathbf{x}|\text{Model 2}) + P(\mathbf{x}|\text{Model 3})) \\ \log P(\text{Model } i|\mathbf{x}) &= \log P(\mathbf{x}|\text{Model } i) - \log (\exp\{\log P(\mathbf{x}|\text{Model 1})\} + \exp\{P(\mathbf{x}|\text{Model 2})\} + \exp\{P(\mathbf{x}|\text{Model 3})\}) \\ \log P(\text{Model } i|\mathbf{x}) &= \log P(\mathbf{x}|\text{Model } i) - \log \left(\sum_{j=1}^3 \exp\{\log P(\mathbf{x}|\text{Model } j)\} \right) \end{aligned} \quad (2.26)$$

For which the second argument can easily be calculated using the *logsumexp()* function from *scipy.special* library.

From the code used, we obtained the following results for the different models log likelihood, and for their probability (**not** in log), all summed in Table 2.1.

	Model 1	Model 2	Model 3
log likelihood	-4.436×10^3	-4.284×10^3	-3.851×10^3
MAP probability	9.143×10^{-255}	1.434×10^{-188}	$1.0 - 9.143 \times 10^{-255} - 1.434 \times 10^{-188}$

Table 2.1: Log Likelihood and MAP estimations of each model

Interpretation What we can deduce from these results is that it is highly likely the data follows a model with each component being Bernoulli distributed with a separate and unknown p_d , compared to Models 1 and 2. This makes sense since our data is diverse and hence make it impossible to follow Model 1.

Listing 2: Code - Q2

```
1 import numpy as np
2 import scipy.special as sp
3 from scipy.special import betaln, logsumexp
4
5 def loglikelihood_model1(Y: np.ndarray, pd = 0.5):
6     N, D = np.shape(Y)
7
8     # p(x given model 1)
9     log_like_model1 = N*D * np.log(pd)
10
11     return log_like_model1
12
13
14 def loglikelihood_model2(Y: np.ndarray):
15     N, D = np.shape(Y)
16     sum_x = np.sum(Y).astype(int)
17
18     log_like_model2 = betaln(1 + sum_x, 1 + (N*D - sum_x))
19
20     return log_like_model2
21
22
23 def loglikelihood_model3(Y: np.ndarray):
24     N, D = np.shape(Y)
25
26     sum_xn = np.zeros((D, 1))
27     log_betas = np.zeros((D, 1))
28
29     for d in range(D):
30         sum_xn[d] = np.sum(Y[:, d]).astype(int)
31         log_betas[d] = betaln(1 + sum_xn[d], N + 1 - sum_xn[d])
32
33     log_like_model3 = np.sum(log_betas)
34
35     return log_like_model3
36
37 def map_prob_model(model_num: int, Y, pd=0.5, log=False):
38
39     if model_num==1:
40         log_like_numerator = loglikelihood_model1(Y, pd)
41     elif model_num==2:
42         log_like_numerator = loglikelihood_model2(Y)
43     else:
44         log_like_numerator = loglikelihood_model3(Y)
45
46     log_like_all = np.array([loglikelihood_model1(Y, pd),
47                             loglikelihood_model2(Y),
```

```
48         loglikelihood_model3(Y)])
49
50     log_map_model = log_like_numerator - logsumexp(log_like_all)
51
52     if not log:
53         map_prob_model = np.exp(log_map_model)
54     else:
55         map_prob_model = log_map_model
56
57     return map_prob_model
58
59 def main():
60     map_model1 = map_prob_model(1, Y)
61     print("The MAP probability of Model 1 is: ", map_model1, "\n")
62     map_model2 = map_prob_model(2, Y)
63     print("The MAP probability of Model 2 is: ", map_model2, "\n")
64     map_model3 = map_prob_model(3, Y)
65     print("The MAP probability of Model 3 is: ", map_model3, "\n")
66
67     like1 = loglikelihood_model1(Y)
68     print("The log likelihood of Model 1 is: ", like1, "\n")
69     like2 = loglikelihood_model2(Y)
70     print("The log likelihood of Model 2 is: ", like2, "\n")
71     like3 = loglikelihood_model3(Y)
72     print("The log likelihood of Model 3 is: ", like3, "\n")
73
74     main()
```

3 | Question 3 - EM for Binary Data

3.1 Question 3.(a)

We consider a mixture of K multivariate Bernoulli distributions. We use the parameters π_1, \dots, π_K to denote the mixing proportions, each comprised between 0 and 1 ($0 \leq \pi_k \leq 1$) and all summing to 1 ($\sum_{k=1}^K \pi_k = 1$).

We define the parameters vectors as $\mathbf{p}_k = (p_{k,1}, \dots, p_{k,D})$, and the matrix P as:

$$\mathbf{P} = [\mathbf{p}_1, \dots, \mathbf{p}_K]^T \quad (3.1)$$

Where each $p_{k,d}$ is defined as $0 \leq p_{k,d} \leq 1$, $k \in \llbracket 0; K \rrbracket$, $d \in \llbracket 0; D \rrbracket$.

Each model is independent and identically distributed (*Assumption 1*), and pixels are independent of each other within each component distribution (*Assumption 2*).

Each of the N images (with index n), assumed to be independent (*Assumption 3*) are defined as:

$$x^{(n)} = (x_1^{(n)}, \dots, x_D^{(n)}) \quad x_d^{(n)} \in \{0, 1\}, \quad d \in \llbracket 0; D \rrbracket, \quad n \in \llbracket 0; N \rrbracket \quad (3.2)$$

Once all these information have been made clear, we can make the following expression, due to [Assumption 2](#):

$$p(x^{(n)}|k) = \prod_{d=1}^D p_{k,d}^{x_d^{(n)}} (1 - p_{k,d})^{1-x_d^{(n)}} \quad (3.3)$$

Leading to the following expression for the likelihood of the mixture element π_k , which is the product between its proportion and its likelihood:

$$p(x^{(n)}|\pi_k) = \pi_k \times p(x^{(n)}|k) \quad (3.4)$$

$$p(x^{(n)}|\pi_k) = \pi_k \prod_{d=1}^D p_{k,d}^{x_d^{(n)}} (1 - p_{k,d})^{1-x_d^{(n)}} \quad (3.5)$$

From that, the probability of x_i given by the set of Multivariate Bernoulli distribution used in the mixture is:

$$p(x^{(n)}|\boldsymbol{\pi}) = \sum_{k=1}^K p(x^{(n)}|\pi_k) \quad (3.6)$$

$$p(x^{(n)}|\boldsymbol{\pi}) = \sum_{k=1}^K \pi_k \prod_{d=1}^D p_{k,d}^{x_d^{(n)}} (1 - p_{k,d})^{1-x_d^{(n)}} \quad (3.7)$$

Thus, for the whole set of images, the likelihood is given by:

$$p(\mathbf{x}|\boldsymbol{\pi}) = p(x^{(1)}, \dots, x^{(N)}|\boldsymbol{\pi}) \quad (3.8)$$

And since [Assumption 3](#):

$$p(\mathbf{x}|\boldsymbol{\pi}) = \prod_{n=1}^N p(x^{(n)}|\boldsymbol{\pi})$$

$$p(\mathbf{x}|\boldsymbol{\pi}) = \prod_{n=1}^N \left(\sum_{k=1}^K \pi_k \prod_{d=1}^D p_{k,d}^{x_d^{(n)}} (1 - p_{k,d})^{1-x_d^{(n)}} \right) \quad (3.9)$$

Taking the log of this expression, we get the log likelihood expression:

$$\begin{aligned} \text{Log Likelihood} &= \log(p(\mathbf{x}|\boldsymbol{\pi})) \\ \text{Log Likelihood} &= \sum_{n=1}^N \log \left(\sum_{k=1}^K \pi_k \prod_{d=1}^D p_{k,d}^{x_d^{(n)}} (1 - p_{k,d})^{1-x_d^{(n)}} \right) \end{aligned} \quad (3.10)$$

For simplicity, we can introduce some matrix notations:

$$\mathbf{X} = [x^{(1)}, \dots, x^{(N)}]^T \quad (3.11)$$

which will be a matrix of size $N \times D$.

We also notice the following:

$$\begin{aligned} p_{k,d}^{x_d^{(n)}} (1 - p_{k,d})^{1-x_d^{(n)}} &= \exp \left\{ \log p_{k,d}^{x_d^{(n)}} (1 - p_{k,d})^{1-x_d^{(n)}} \right\} \\ &= \exp \left\{ x_d^{(n)} \log p_{k,d} + (1 - x_d^{(n)}) \log(1 - p_{k,d}) \right\} \end{aligned}$$

Thus, replacing in 3.7, we get:

$$\begin{aligned} p(x^{(n)}|\boldsymbol{\pi}) &= \sum_{k=1}^K \pi_k \prod_{d=1}^D p_{k,d}^{x_d^{(n)}} (1 - p_{k,d})^{1-x_d^{(n)}} \\ &= \sum_{k=1}^K \pi_k \prod_{d=1}^D \exp \left\{ x_d^{(n)} \log p_{k,d} + (1 - x_d^{(n)}) \log(1 - p_{k,d}) \right\} \\ p(x^{(n)}|\boldsymbol{\pi}) &= \sum_{k=1}^K \pi_k \exp \left\{ \sum_{d=1}^D \left(x_d^{(n)} \log p_{k,d} + (1 - x_d^{(n)}) \log(1 - p_{k,d}) \right) \right\} \end{aligned} \quad (3.12)$$

And using now matrices \mathbf{X} and \mathbf{P} of dimensions $N \times K$ and $K \times D$ respectively:

$$p(x^{(n)}|\boldsymbol{\pi}) = \sum_{k=1}^K \pi_k \exp \left\{ (x^{(n)})^T \log(\mathbf{p}_k) + (1 - x^{(n)})^T \log(1 - \mathbf{p}_k) \right\} \quad (3.13)$$

Generalizing for the log likelihood ¹²:

$$\log p(\mathbf{x}|\boldsymbol{\pi}) = \sum_{n=1}^N \log \sum_{k=1}^K \pi_k \exp \left\{ (x^{(n)})^T \log(\mathbf{p}_k) + (1 - x^{(n)})^T \log(1 - \mathbf{p}_k) \right\} \quad (3.14)$$

¹(since we take the log, it will be a sum of the logs instead of the product of the elements)

² x^n are vectors, although not written bold

3.2 Question 3.(b)

We would like to find an expression for the responsibility of the mixture component k , denoted as $r_{n,k}$, for data vector $\mathbf{x}^{(n)}$. Using Bayes rule:

$$r_{n,k} = P(s^{(n)} = k | \mathbf{x}^{(n)}, \boldsymbol{\pi}, \mathbf{P}) = \frac{P(\mathbf{x}^{(n)} | s^{(n)} = k, \boldsymbol{\pi}, \mathbf{P}) P(s^{(n)} = k | \boldsymbol{\pi}, \mathbf{P})}{P(\mathbf{x}^{(n)} | \boldsymbol{\pi}, \mathbf{P})} \quad (3.15)$$

We notice that, since $s^{(n)}$ does not depend on \mathbf{P} :

$$P(s^{(n)} = k | \boldsymbol{\pi}, \mathbf{P}) = P(s^{(n)} = k | \boldsymbol{\pi}) = \pi_k$$

And using the sum and product rules of probabilities, we get:

$$P(\mathbf{x}^{(n)} | \boldsymbol{\pi}, \mathbf{P}) = \sum_{i=1}^K P(\mathbf{x}^{(n)}, s^{(n)} = i | \boldsymbol{\pi}, \mathbf{P}) = \sum_{i=1}^K P(\mathbf{x}^{(n)} | s^{(n)} = i, \boldsymbol{\pi}, \mathbf{P}) \times P(s^{(n)} = i | \boldsymbol{\pi}, \mathbf{P})$$

Therefore, replacing in our expression, we get for the responsibility of mixture component k for data vector $\mathbf{x}^{(n)}$:

$$\begin{aligned} r_{n,k} &= \frac{P(\mathbf{x}^{(n)} | s^{(n)} = k, \boldsymbol{\pi}, \mathbf{P}) \times P(s^{(n)} = k | \boldsymbol{\pi}, \mathbf{P})}{\sum_{i=1}^K P(\mathbf{x}^{(n)} | s^{(n)} = i, \boldsymbol{\pi}, \mathbf{P}) \times P(s^{(n)} = i | \boldsymbol{\pi}, \mathbf{P})} \\ r_{n,k} &= \frac{\pi_k P(\mathbf{x}^{(n)} | s^{(n)} = k, \boldsymbol{\pi}, \mathbf{P})}{\sum_{i=1}^K \pi_i P(\mathbf{x}^{(n)} | s^{(n)} = i, \boldsymbol{\pi}, \mathbf{P})} \end{aligned} \quad (3.16)$$

And since, using the E-step of the EM algorithm, we have:

$$P(\mathbf{x}^{(n)} | s^{(n)} = k, \boldsymbol{\pi}, \mathbf{P}) = \prod_{d=1}^D p_{k,d}^{x_d^{(n)}} (1 - p_{k,d})^{1-x_d^{(n)}}$$

Then:

$$r_{n,k} = \frac{\pi_k \prod_{d=1}^D p_{k,d}^{x_d^{(n)}} (1 - p_{k,d})^{1-x_d^{(n)}}}{\sum_{i=1}^K \pi_i \prod_{d=1}^D p_{i,d}^{x_d^{(n)}} (1 - p_{i,d})^{1-x_d^{(n)}}} \quad (3.17)$$

Taking the log and naming R_k the expression in the numerator, we get:

$$\begin{aligned} R_k &= \pi_k \prod_{d=1}^D p_{k,d}^{x_d^{(n)}} (1 - p_{k,d})^{1-x_d^{(n)}} \\ \log R_k &= \log \pi_k + \sum_{d=1}^D (x_d^{(n)} \log p_{k,d} + (1 - x_d^{(n)}) \log(1 - p_{k,d})) \end{aligned}$$

And:

$$\log r_{n,k} = \log R_k - \log \left(\sum_{i=1}^K R_i \right) \quad (3.18)$$

3.3 Question 3.(c)

We would like to find the maximizing parameters for the expected log-joint with respect to parameters $\boldsymbol{\pi}$ and \mathbf{P} . It can be expressed as:

$$\left\langle \sum_{n=1}^N \log P(\mathbf{x}^{(n)}, s^{(n)} | \boldsymbol{\pi}, \mathbf{P}) \right\rangle_{q(s^{(n)})} = \sum_{n=1}^N q(s^{(n)}) \log P(\mathbf{x}^{(n)}, s^{(n)} | \boldsymbol{\pi}, \mathbf{P}) \quad (3.19)$$

Using conditional probabilities:

$$\log P(\mathbf{x}^{(n)}, s^{(n)} | \boldsymbol{\pi}, \mathbf{P}) = \log P(\mathbf{x}^{(n)} | s^{(n)}, \boldsymbol{\pi}, \mathbf{P}) + \log P(s^{(n)} | \boldsymbol{\pi}, \mathbf{P})$$

for which we already found the expression before (in question (b)) being:

$$\log P(\mathbf{x}^{(n)}, s^{(n)} | \boldsymbol{\pi}, \mathbf{P}) = \log \boldsymbol{\pi} + \sum_{d=1}^D (x_d^{(n)} \log p_{k,d} + (1 - x_d^{(n)}) \log(1 - p_{k,d})) \quad (3.20)$$

Or better in matrix and vector form:

$$\log P(\mathbf{x}^{(n)}, s^{(n)} | \boldsymbol{\pi}, \mathbf{P}) = \log \boldsymbol{\pi} + \sum_{d=1}^D \log(\mathbf{P})^T \mathbf{x}^{(n)} + \log(1 - \mathbf{P})^T (1 - \mathbf{x}^{(n)})$$

From before, we know the corresponding E-step in our case is:

$$q(s^{(n)}) = \mathbf{r}_n = [r_{n,1}, \dots, r_{n,K}]^T$$

Resulting in the expression of E:

$$E = \sum_{n=1}^N \mathbf{r}_n \left(\log \boldsymbol{\pi} + \sum_{d=1}^D \mathbf{x}^{(n)} * \log(\mathbf{P})^T + (1 - \mathbf{x}^{(n)}) \log(1 - \mathbf{P})^T \right) \quad (3.21)$$

Maximizing the E-step:

$$\operatorname{argmax}_{\boldsymbol{\pi}, \mathbf{P}} \left\langle \sum_{n=1}^N \log P(\mathbf{x}^{(n)}, s^{(n)} | \boldsymbol{\pi}, \mathbf{P}) \right\rangle_{q(s^{(n)})} \quad (3.22)$$

Is equivalent to take the derivative of the expression of E and differentiate it with respect to $\boldsymbol{\pi}$ and \mathbf{P} and set the derivatives to 0.

Maximizing \mathbf{P} , by finding the optimal $\hat{p}_{k,d}$:

$$\begin{aligned} \frac{\partial E}{\partial p_{k,d}} &= \frac{\partial}{\partial p_{k,d}} \sum_{n=1}^N \mathbf{r}_n \left(\log \boldsymbol{\pi} + \sum_{d=1}^D \mathbf{x}^{(n)} \log(\mathbf{P})^T + (1 - \mathbf{x}^{(n)}) \log(1 - \mathbf{P})^T \right) \\ &= \sum_{n=1}^N r_{n,k} \frac{\partial}{\partial p_{k,d}} \left(\log(p_{k,d}) x_d^{(n)} + \log(1 - p_{k,d}) (1 - x_d^{(n)}) \right) \\ &= \sum_{n=1}^N r_{n,k} \left(\frac{x_d^{(n)}}{p_{k,d}} - \frac{(1 - x_d^{(n)})}{1 - p_{k,d}} \right) \end{aligned} \quad (3.23)$$

And setting this equation to 0:

$$\frac{\partial E}{\partial p_{k,d}} = 0 \quad (3.24)$$

$$\begin{aligned} \sum_{n=1}^N r_{n,k} \left(\frac{x_d^{(n)}}{\hat{p}_{k,d}} + \frac{-(1-x_d^{(n)})}{1-\hat{p}_{k,d}} \right) &= 0 \\ \frac{(1-\hat{p}_{k,d}) \sum_{n=1}^N r_{n,k} x_d^{(n)} - \hat{p}_{k,d} \sum_{n=1}^N r_{n,k} (1-x_d^{(n)})}{\hat{p}_{k,d}(1-\hat{p}_{k,d})} &= 0 \\ \sum_{n=1}^N r_{n,k} x_d^{(n)} - \hat{p}_{k,d} \sum_{n=1}^N r_{n,k} x_d^{(n)} + \hat{p}_{k,d} \sum_{n=1}^N r_{n,k} x_d^{(n)} - \hat{p}_{k,d} \sum_{n=1}^N r_{n,k} &= 0 \\ \sum_{n=1}^N r_{n,k} x_d^{(n)} - \hat{p}_{k,d} \sum_{n=1}^N r_{n,k} &= 0 \\ \hat{p}_{k,d} &= \frac{\sum_{n=1}^N r_{n,k} x_d^{(n)}}{\sum_{n=1}^N r_{n,k}} \end{aligned} \quad (3.25)$$

Doing the same for π to find the optimal $\hat{\pi}_k$:

$$\frac{\partial E}{\partial \pi_k} = 0$$

However, this equation result leads to the sum over n of $r_{n,k}$ being equal to 0, because:

$$\frac{\partial E}{\partial \pi_k} = \sum_{n=1}^N \frac{r_{n,k}}{\pi_k}$$

This doesn't solve our problem. We must redefine this problem as an optimization problem where we want to maximize E subject to the constraint $\sum_k \pi_k = 1$.

Using the Lagrangian multiplier λ to enforce normalization:

$$E_{Lagrangian} = E + \lambda \left(1 - \sum_{k=1}^K \pi_k \right) = 0 \quad (3.26)$$

And finding the conditions for stationarity:

$$\begin{aligned} \frac{\partial E_{Lagrangian}}{\partial \pi_k} = 0 &\iff \sum_{n=1}^N \frac{r_{n,k}}{\hat{\pi}_k} - \lambda = 0 \\ \sum_{n=1}^N r_{n,k} &= \lambda \hat{\pi}_k \\ \hat{\pi}_k &= \frac{\sum_{n=1}^N r_{n,k}}{\lambda} \end{aligned} \quad (3.27)$$

And:

$$\frac{\partial E_{Lagrangian}}{\partial \lambda} = 0$$

Leads to:

$$\sum_{k=1}^K \pi_k = 1 \quad (3.28)$$

To find the parameter λ , we now replace in 3.28:

$$\sum_{k=1}^K \left(\frac{\sum_{n=1}^N r_{n,k}}{\lambda} \right) = 1$$
$$\lambda = \sum_{k=1}^K \sum_{n=1}^N r_{n,k} \quad (3.29)$$

Therefore, the final expression for the optimal $\hat{\pi}_k$ is:

$$\hat{\pi}_k = \frac{\sum_{n=1}^N r_{n,k}}{\sum_{k=1}^K \sum_{n=1}^N r_{n,k}} = \frac{\sum_{n=1}^N r_{n,k}}{\sum_{n=1}^N \sum_{k=1}^K r_{n,k}}$$

And since the sum over k of all the responsibilities $r_{n,k}$ is 1, then:

$$\hat{\pi}_k = \frac{\sum_{n=1}^N r_{n,k}}{\sum_{n=1}^N 1}$$
$$\hat{\pi}_k = \frac{1}{N} \sum_{n=1}^N r_{n,k} \quad (3.30)$$

3.4 Question 3.(d)

After implementing the EM algorithm for a mixture of K multivariate Bernoullis, the obtained results are displayed in the figures below (Figures 3.1 to 3.5), where the code was executed for $K \in \{2, 3, 4, 7, 10\}$. The initial value of $\boldsymbol{\pi}$ has been chosen as:

$$\boldsymbol{\pi} = \left[\frac{1}{\sum_k k}, \dots, \frac{K}{\sum_k k} \right]$$

and for \mathbf{P} were for each $p_{k,d}$ the mean values of $x_d^{(n)}$ (the same values were given on dimension k).

We can see that for $K = 10$, the algorithm converges faster to a better value. This makes sense as there are 9 different single digits that can be drawn. The model would therefore find approximately a category for each, and one last category for the ones that were badly drawn.

A notable issue about it is the fact the optimization of the EM parameters only happened after around 5 to 10 iterations. We can suppose the model's performance depends on the initial values.

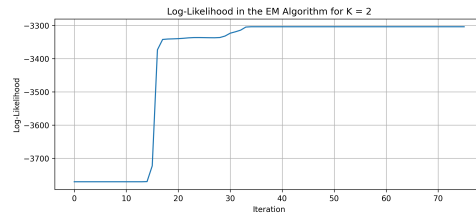


Figure 3.1: K=2

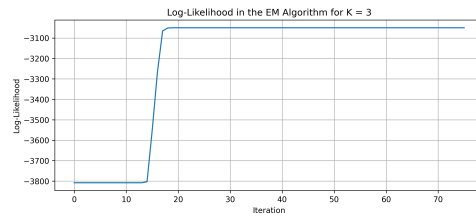


Figure 3.2: K=3

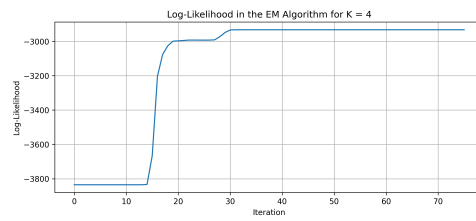


Figure 3.3: K=4

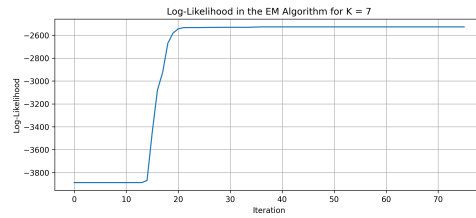


Figure 3.4: K=7

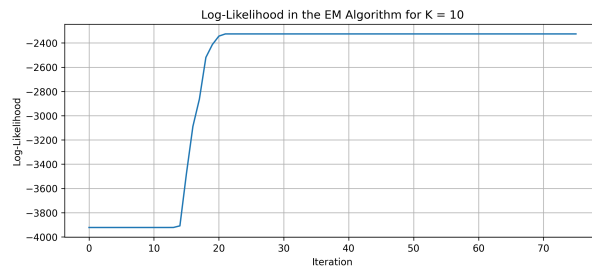


Figure 3.5: K=10

```
1 import numpy as np
2 from scipy.special import logsumexp
3 import matplotlib.pyplot as plt
4
5 # Defining the algorithm function, which takes K (number of
6 # mixture components), the matrix X (containing dataset) and
7 # the maximum iterations to run
8
9 def EM_algorithm(K: int,
10                 X: np.ndarray,
11                 n_iter: int=300,
12                 epsilon: float=1e-12,
13                 min_break = 75
14                 ):
15     """
16     Main function fo executing the EM algorithm
17
18     Inputs:
19     - K: int, the number of mixture components
20     - X: int, the matrix containing the dataset
21     - n_iter: int, the maximum number of iterations
22     - epsilon: float, the precision after which we can stop
23       the algorithm from running.
24
25     Outputs:
26     - pi_k: array, size (K,1) containing all mixing proportions
27     - p_kd: array, size (K,D) containing all p_kds (equivalent to
28       matrix P in problem wording)
29     - log_likelihood: list, containing all the log likelihood
30       updates of the model.
31     """
32
33     # Finding dimensions of initial dataset
34     N, D = np.shape(X)
35
36     # Defining initial pi_k as an array
37     # with increasing values
38     pi_k = np.arange(1,K+1)
39     sum_pi = np.sum(pi_k)
40     pi_k = (1/sum_pi)*pi_k
41
42     # Initializing P's values as mean of
43     # element d in all images
44     p_kd = np.zeros((K, D))
45     for d in range(D):
46         p_kd[:, d] = (1/N)*X[:, d].sum()
47
48     # Initializing the final
```

```

49 log_likelihood = []
50
51 # Checking for any 0 or 1 in the values of pi
52 # and P, and if so, replacing them by very small
53 # numbers. It avoids numerical instability from
54 # Dividing by a 0 or log(1)=0
55 for k in range(K):
56     if pi_k[k] < 1e-10:
57         pi_k[k] = 1e-10
58     elif pi_k[k] > (1- 1e-10):
59         pi_k[k] = 1- 1e-10
60
61     for d in range(D):
62         if p_kd[k,d] < 1e-10:
63             p_kd[k,d] = 1e-10
64         elif p_kd[k,d] > 1-1e-10:
65             p_kd[k,d] = 1-1e-10
66
67 # Loop of the EM algorithm, which will stop when the
68 # max number of iterations will be reached OR when the
69 # updates won't have a difference superior to epsilon
70 for i in range(n_iter):
71
72     # Calculating the responsibilities (E-step)
73     r_nk = E_step(X, pi_k, p_kd, K)
74
75     # Calculating the new pi and P values (M-Step)
76     pi_k, p_kd = M_step(X, r_nk)
77
78     # Checking for invalid values to avoid numerical
79     # instability
80     for k in range(K):
81         if pi_k[k] < 1e-10:
82             pi_k[k] = 1e-10
83         elif pi_k[k] > (1- 1e-10):
84             pi_k[k] = 1- 1e-10
85         for d in range(D):
86             if p_kd[k,d] < 1e-10:
87                 p_kd[k,d] = 1e-10
88             elif p_kd[k,d] > 1-1e-10:
89                 p_kd[k,d] = 1-1e-10
90
91     # Calculating the new log likelihood, given the calculated
92     # optimal parameters and storing it in the list
93     log_like = log_likelihood_EM(X, pi_k, p_kd, r_nk)
94     log_likelihood.append(log_like)
95
96     # Breaking the loop if the updates don't bring any change
97     if i > min_break and abs(log_likelihood[-1] - log_likelihood[-2]) < epsilon:

```

```

98         break
99
100     return pi_k, p_kd, log_likelihood
101
102 def E_step(X: np.ndarray, pi_k, p_kd, K):
103     """
104     Calculating the E-step of the algorithm, using the equations found earlier
105
106     Inputs:
107     - X: np.ndarray, shape (N,D), contains initial data
108     - pi_k: list, shape(K,), contains current pi values for each mixture
109     - p_kd: array, shape(K, D), contains current probabilities
110     - K: int, number of models in the mixture
111
112     Outputs:
113     - r_nk: np.ndarray, shape(N,K), contains responsibilities of each model
114     """
115     # Defining the values of N, being number of rows in dataset
116     N, _ = np.shape(X)
117
118     # Initializing the responsibilities matrix
119     log_r_nk = np.zeros((N, K))
120
121     # Reuniting the different rows of pi_k
122     # as a single row
123     log_pi_k = np.log(pi_k).ravel()
124
125     # Calculating the logs of each responsibility
126     log_r_nk = log_pi_k + (X @ np.log(p_kd.T)) + ((1 - X) @ np.log(1 - p_kd.T))
127
128     # Normalizing the logs of responsibilities
129     log_r_nk = log_r_nk - logsumexp(log_r_nk, axis=1, keepdims=True)
130
131     # Returning to the non-log domain, by calculating the exponential
132     # of the responsibilities
133     r_nk = np.exp(log_r_nk)
134
135     return r_nk
136
137 def M_step(X: np.ndarray, r_nk):
138     """
139     Calculates the M-step of the EM algorithm, using the equations defined in report
140
141     Inputs:
142     - X: np.ndarray, shape (N,D), contains initial data
143     - r_nk: array, shape (N, K), responsibilities calculated of each model on
144           each data point
145
146     Outputs:
147     - pi_k: list, shape (K,), contains current pi values for each mixture

```

```

147     - P: array, shape (K, D), contains current probabilities
148     """
149
150     N, D = np.shape(X)
151     K = np.shape(r_nk)[1]
152
153     # Verifying the validity of the values
154     # of the responsibilities
155     for n in range(N):
156         for k in range(K):
157             if r_nk[n,k] < 1e-10:
158                 r_nk[n, k] = 1e-10
159             elif r_nk[n,k] > (1-1e-10):
160                 r_nk[n,k] = 1-1e-10
161
162     # Calculating the updates of the parameters
163     pi_k = (1/N)*np.sum(r_nk, axis=0)
164     log_p_kd = np.log(r_nk.T @ X) - np.log(r_nk.sum(axis=0)[:, np.newaxis])
165     p_kd = np.exp(log_p_kd)
166
167     # Verifying the validity of the values
168     # of the newly calculated parameters (in pi and P)
169     for k in range(K):
170         if pi_k[k] < 1e-10:
171             pi_k[k] = 1e-10
172         elif pi_k[k] > (1- 1e-10):
173             pi_k[k] = 1- 1e-10
174         for d in range(D):
175             if p_kd[k,d] < 1e-10:
176                 p_kd[k,d] = 1e-10
177             elif p_kd[k,d] > 1-1e-10:
178                 p_kd[k,d] = 1-1e-10
179
180     return pi_k, p_kd
181
182
183 def log_likelihood_EM(X, pi_k, p_kd, r_nk):
184     """
185     Calculating the log likelihood of the data being given the parameters
186     of the mixture model
187
188     Inputs:
189     - X: np.ndarray, shape (N,D), contains initial data
190     - pi_k: list, shape (K,), contains current pi values for each mixture
191     - p_kd: array, shape (K, D), contains current probabilities
192     - r_nk: array, shape (N, K), responsibilities calculated of each model on
193           each data point
194
195     Outputs:
196     - log_likelihood: float, log likelihood probability calculated

```

```

196     """
197
198     # Calculating N, D, K and initializing the value of the log likelihood
199     N, D = np.shape(X)
200     K = np.shape(pi_k)[0]
201     log_likelihood = 0
202
203     # Looping through the r_nk dimensions, since there will be as many elements to add
204     # as there is responsibilities
205     for n in range(N):
206         for k in range(K):
207
208             # Calculating the log of element k in pi matrix
209             log_pi = np.log(pi_k[k])
210
211             # Calculating  $P(x | \text{vert } s=k, \text{pi}, P)$ , the main element in the calculation
212             log_px_given_k = np.sum(X[n] * np.log(p_kd[k])
213                                     + (1 - X[n]) * np.log(1 - p_kd[k]))
214
215             # Weighting the log of probabilities by the responsibilities, and
216             # adding them to the log likelihood.
217             log_rnk_contribution = r_nk[n,k] * (log_pi + log_px_given_k)
218             log_likelihood += log_rnk_contribution
219
220     return log_likelihood
221
222
223 def plot_EM(log_likelihoods, pi_k, p_kd, K, save_like=False, save_prob=False):
224     """
225     Plots the EM results obtained
226     """
227
228     # Plotting the obtained log likelihoods as a function of the iterations
229     plt.figure(figsize=(10,4))
230     plt.plot(log_likelihoods)
231     plt.xlabel("Iteration")
232     plt.ylabel("Log-Likelihood")
233     plt.title(f"Log-Likelihood in the EM Algorithm for K = {K}")
234     plt.grid()
235     if save_like:
236         plt.savefig(f'Q3/Loglikelihood_{K}.png',
237                   format="png",
238                   dpi=300,
239                   bbox_inches="tight")
240     plt.show()
241
242
243     # Plotting the mixture components probabilities as heat maps
244     K, D = p_kd.shape

```

```
245     image_size = int(np.sqrt(D)/2)
246     figs, axs = plt.subplots(1, K, figsize=(4*K, 4))
247     figs.suptitle(f"Mixture components probabilities for K = {K}", fontsize=14)
248     for i in range(K):
249         ax = axs[i] if K > 1 else axs
250         ax.imshow(p_kd[i].reshape((image_size, image_size)), cmap='viridis')
251         ax.set_title(f"pi: {pi_k[i]:.2f}", fontsize=10)
252         ax.axis('off')
253     if save_prob:
254         plt.savefig(f'Q3/prob_{K}.png', format="png", dpi=300, bbox_inches="tight")
255     plt.show()
256
257
258 def main():
259     Ks: list = [2, 3, 4, 7, 10]
260
261     for K in Ks:
262         print("=== EM Algorithm for K =", K, " ===")
263         pi_k, p_kd, log_likelihoods = EM_algorithm(K, Y, n_iter=75, epsilon=1e-6)
264         print("For K=", K, ", \n pi_k = ", pi_k, "\n p_kd = ", p_kd) print("\n and log_likelihoods = ")
265         plot_EM(log_likelihoods, pi_k, p_kd, K)
266
267
268 main()
```

3.5 Question 3.(e)

In this section, we used the code from the previous questions and slightly modifying it (see the code following the figures below) to make it generate random initial parameters. The different initial random probabilities were normalized after being generated to still be valid probabilities.

In Figure ?? are displayed the initial random $p_{k,d}$ used in the EM algorithm. We can see there aren't any relevant pattern. Then, by running the algorithm on such initial values, we obtain the final parameters learned by the EM algorithm and displayed in Figure 3.11. The resulting parameters can be good for low dimensions (*e.g.* 3.6), but for larger dimensions, rapidly, it has learned parameters that do not look very good (with extreme probabilities, *i.e.* either 0 or 1). Comparing them to that of our fairly smart proposal (the data-driven one, shown in Figure 3.17) and by comparing both, it seems each category the data-driven look like a real handwritten digit.

We therefore conclude we got different solutions depending on the starting point. This algorithm therefore has some flaws, and the main flaw is its over-dependence on the data. By choosing the wrong data, the model won't be able to classify well the different categories of handwritten digits. From looking at the pictures, we also see it fails at finding clusters, since the handwritten digits are all similar in the different categories. Thus, overall, it would perform pretty badly in a real life scenario.

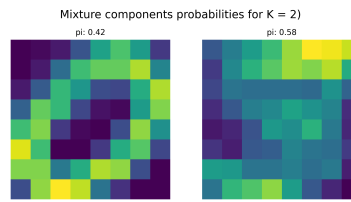


Figure 3.6: $K=2$

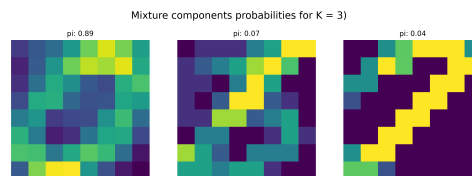


Figure 3.7: $K=3$

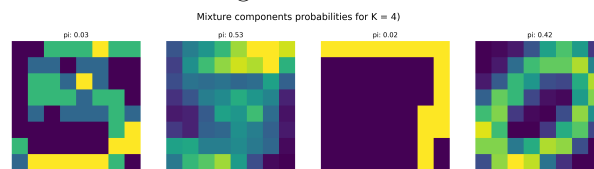


Figure 3.8: $K=4$

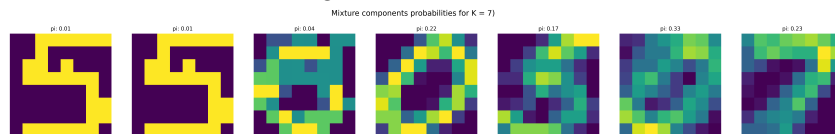


Figure 3.9: $K=7$

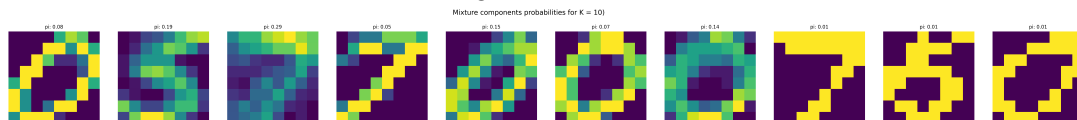


Figure 3.10: $K=10$

Figure 3.11: Learned Parameters from Random Initial Parameters

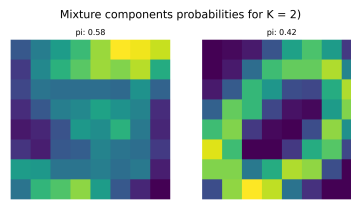


Figure 3.12: $K=2$

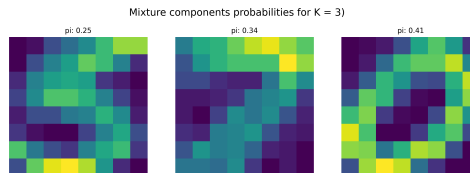


Figure 3.13: $K=3$

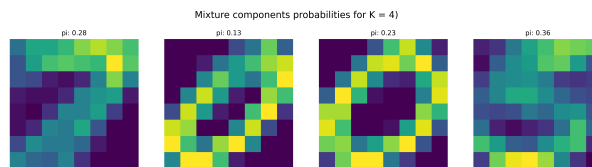


Figure 3.14: $K=4$

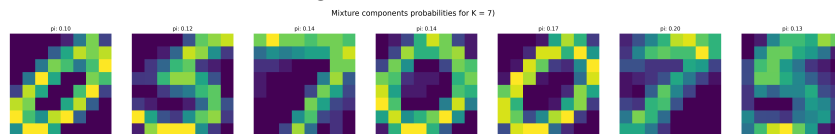


Figure 3.15: $K=7$

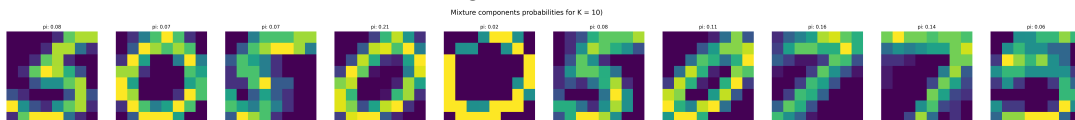


Figure 3.16: $K=10$

Figure 3.17: Parameters Learned by the EM Algorithm for Data-Driven Initialization

4 | Question 5 - Decrypting Messages with MCMC

We are given an encrypted passage of English text. The mapping of this encryption is one-to-one, in a sense that all encrypted symbols are assigned to a unique other different symbol. The English text, composed of s_i symbols as $s_1 s_2 \dots s_n$, is modelled as a first-order Markov Chain:

$$p(s_1 s_2 \dots s_n) = p(s_1) \prod_{i=2}^n p(s_i | s_{i-1}) \quad (4.1)$$

4.1 Question 5.(a)

We define the transition probabilities in the studied text (*i.e.* "War and Peace", by Leo Tolstoy), as $p(s_i = \alpha | s_{i-1} = \beta) = \psi(\alpha, \beta)$ and the stationary distributions of the symbols as $\lim_{i \rightarrow \infty} p(s_i = \gamma) = \phi(\gamma)$. We assume the first letter of the encrypted text is sampled from the stationary distributions.

The formulae for the ML estimates of $p(s_i = \alpha | s_{i-1} = \beta) = \psi(\alpha, \beta)$ are given by:

$$\psi_{ML}(\alpha, \beta) = \frac{\text{counts of pair } (\alpha, \beta) \text{ as } \beta\alpha \text{ in the text}}{\text{counts of symbol } \beta \text{ in the text}} \quad (4.2)$$

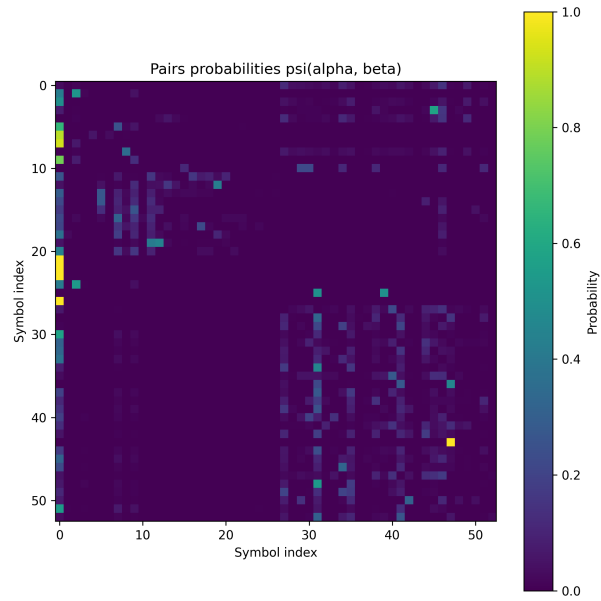
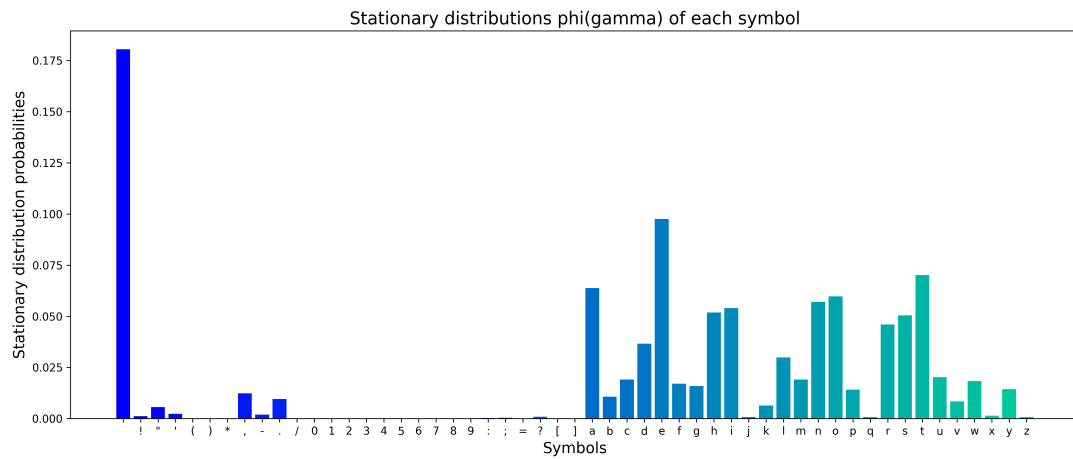
The stationary probabilities were calculated in the code by calculating, using the `np.linalg.eig()` function on the transpose of the transition probability matrix, finding the eigenvector corresponding to the eigenvalue with the value the closest to 1 and normalizing the values in that eigenvector by dividing them by their entire sum.

In Table 4.1, a few of the transition probability matrix are displayed. Since the table is too large, the csv file containing the whole matrix is joint to the report. In addition, Figure 4.1 shows the heat map plot corresponding to the matrix, which is clearer.

In Table 4.3 are displayed all the entries of the stationary distribution. In addition, the bar plot shown in Figure 4.2 shows the stationary distribution probabilities for each of the corresponding symbols.

	a	b	c	d	e	f	g	h
a	3.401e-05	1.681e-02	3.410e-02	5.500e-02	8.405e-04	8.104e-03	1.709e-02	1.526e-03
b	9.254e-02	6.608e-03	8.657e-05	4.617e-04	3.288e-01	0.000e+00	0.000e+00	2.886e-05
c	1.091e-01	0.000e+00	1.702e-02	1.460e-04	2.132e-01	0.000e+00	0.000e+00	1.883e-01
d	2.083e-02	1.014e-04	8.454e-05	1.102e-02	1.163e-01	7.186e-04	3.669e-03	4.227e-04
e	4.231e-02	8.088e-04	1.725e-02	9.085e-02	2.564e-02	1.005e-02	6.835e-03	2.265e-03
f	7.022e-02	4.918e-04	0.000e+00	0.000e+00	8.418e-02	5.163e-02	0.000e+00	5.465e-05
g	6.606e-02	0.000e+00	0.000e+00	8.183e-04	1.146e-01	0.000e+00	8.670e-03	1.161e-01
h	1.657e-01	3.703e-04	3.703e-04	3.584e-04	4.496e-01	4.062e-04	5.973e-06	3.584e-05

Table 4.1: A few transition probabilities displayed

Figure 4.1: Heat Map of the Transition Probabilities for each probability $\psi(\alpha, \beta)$ Figure 4.2: Stationary distributions figure, showing the $\phi(\gamma)$ probability for each corresponding symbol.

Symbol	Stationary Probability
	1.805e-01
!	1.216e-03
"	5.565e-03
,	2.331e-03
(2.065e-04
)	2.065e-04
*	8.917e-05
-	1.880e-03
.	9.561e-03
/	2.787e-06
0	5.264e-05
1	1.208e-04
2	4.459e-05
3	1.827e-05
4	7.122e-06
5	1.579e-05
6	1.641e-05
7	1.177e-05
8	5.945e-05
9	9.908e-06
:	3.118e-04
;	3.542e-04
=	6.193e-07
?	9.707e-04
[6.193e-07
]	6.193e-07
a	6.373e-02
b	1.073e-02
c	1.908e-02
d	3.663e-02
e	9.762e-02
f	1.700e-02
g	1.589e-02
h	5.184e-02
i	5.395e-02
j	7.973e-04
k	6.328e-03
l	2.989e-02
m	1.909e-02
n	5.703e-02
o	5.973e-02
p	1.410e-02
q	7.218e-04
r	4.596e-02
s	5.044e-02
t	7.012e-02
u	2.026e-02
v	8.386e-03
w	1.834e-02
x	1.357e-03
y	1.433e-02
z	7.391e-04

Table 4.2: Stationary Probabilities

4.2 Question 5.(b)

We represent the mapping of s onto its corresponding encoding as $\sigma(s)$. We define each state of the MCMC sampler as a permutation between two mapping variables, meaning we simply exchange randomly the encoding of two distinct variables to their corresponding encoding symbols.

Assuming a uniform prior distribution over the permutations, we ask ourselves if the latent variables $\sigma(s)$ for different symbols s independent. The answer is no: the latent variables are not independent because assigning a latent variable to a symbol s reduces by one choice the possible assignment of the next latent variable to the next symbol s .

Let $e_1 \dots e_n$ be an encrypted English text. We recall the equation (2.1), $p(s_1 s_2 \dots s_n) = p(s_1) \prod_{i=2}^n p(s_i | s_{i-1})$. Then relating them to the encoding function, the joint probability of $e_1 \dots e_n$ given σ is:

$$e_i = \sigma(s_i) \iff s_i = \sigma^{-1}(e_i)$$

And thus:

$$p(e_1, \dots, e_n, | \sigma) = \phi(\sigma^{-1}(e_1)) \prod_{i=2}^n p(\sigma^{-1}(e_i) | \sigma^{-1}(e_{i-1})) \quad (4.3)$$

4.3 Question 5.(c)

We use the Metropolis-Hastings algorithm with the proposal given by choosing two symbols s and s^2 at random and swapping the corresponding encrypted symbols.

The proposal probability $S(\sigma \rightarrow \sigma')$ depends on the permutations σ and σ' since we swap the encoding of two symbols in the mapping. Therefore, the proposal probability is the probability of choosing those two symbols simultaneously, given by:

$$S(\sigma \rightarrow \sigma') = \frac{1}{\binom{n}{2}} \quad (4.4)$$

where n is the number of symbols. In our case, we get:

$$S(\sigma \rightarrow \sigma') = \frac{1}{\binom{53}{2}} = \frac{2! \times 51!}{53!} = \frac{2}{53 \times 52}$$

$$S(\sigma \rightarrow \sigma') = \frac{1}{1378} \quad (4.5)$$

The MH acceptance probability, which depends on the current mapping (while the proposal probability did not) for a given proposal is given by:

$$A(\sigma' \rightarrow \sigma | e_1, \dots, e_n) = \min \left(1, \frac{S(\sigma \rightarrow \sigma') p(\sigma' | e_1, \dots, e_n)}{S(\sigma' \rightarrow \sigma) p(\sigma | e_1, \dots, e_n)} \right)$$

We notice $S(\sigma' \rightarrow \sigma) = S(\sigma \rightarrow \sigma')$. Thus, we are left with:

$$A(\sigma \rightarrow \sigma' | e_1, \dots, e_n) = \min \left(1, \frac{p(\sigma' | e_1, \dots, e_n)}{p(\sigma | e_1, \dots, e_n)} \right) \quad (4.6)$$

Recalling Bayes' Theorem:

$$p(\sigma | e_1, \dots, e_n) = \frac{p(e_1, \dots, e_n | \sigma) p(\sigma)}{p(e_1, \dots, e_n)}$$

And similarly for σ' :

$$p(\sigma'|e_1, \dots, e_n) = \frac{p(e_1, \dots, e_n|\sigma')p(e_1, \dots, e_n)}{p(\sigma')} \iff p(e_1, \dots, e_n) = \frac{p(\sigma')p(\sigma'|e_1, \dots, e_n)}{p(e_1, \dots, e_n|\sigma')}$$

Combining the two expressions (by replacing $p(e_1, \dots, e_n)$):

$$p(\sigma|e_1, \dots, e_n) = \frac{p(e_1, \dots, e_n|\sigma)}{p(\sigma)} \times \frac{p(\sigma')p(\sigma'|e_1, \dots, e_n)}{p(e_1, \dots, e_n|\sigma')} \quad (4.7)$$

And reinjecting them into the expression of A (equation (2.6)):

$$\begin{aligned} A(\sigma \rightarrow \sigma'|e_1, \dots, e_n) &= \min \left(1, \frac{p(\sigma'|e_1, \dots, e_n)}{\frac{p(e_1, \dots, e_n|\sigma)}{p(\sigma)} \times \frac{p(\sigma')p(\sigma'|e_1, \dots, e_n)}{p(e_1, \dots, e_n|\sigma')}} \right) \\ &= \min \left(1, \frac{1}{\frac{p(e_1, \dots, e_n|\sigma)}{p(\sigma)} \times \frac{p(\sigma')}{p(e_1, \dots, e_n|\sigma')}} \right) \end{aligned}$$

Leading to:

$$A(\sigma \rightarrow \sigma'|e_1, \dots, e_n) = \min \left(1, \frac{p(e_1, \dots, e_n|\sigma')p(\sigma)}{p(\sigma')p(e_1, \dots, e_n|\sigma)} \right)$$

And since we assumed a uniform prior distribution over the permutations, then $p(\sigma) = p(\sigma')$ and therefore, the expression of the MH acceptance probability simplifies to:

$$A(\sigma \rightarrow \sigma'|e_1, \dots, e_n) = \min \left(1, \frac{p(e_1, \dots, e_n|\sigma')}{p(e_1, \dots, e_n|\sigma)} \right) \quad (4.8)$$

4.4 Question 5.(d)

This section's goal is to implement the MH sampler and running it on the encrypted text. In Table 4.3 are displayed the results every 100 iterations, using a random initial proposal, for the proposal encoding shown in 4.4.

Encoded symbol	Original symbol
9	=
i	
r	-
''''''	:
!	:
f	!
:	?
c	/
u	:
?	,
*	''''''
s	(
z)
p	
6	
0	*
q	1
o	2
a	3
w	4
h	5
=	6
3	7
2	8
g	9
	a
y	b
k	c
l	d
j	e
)	f
x	g
b	h
7	i
v	j
1	k
/	l
.	m
	n
e	o
5	p
-	q
t	r
	s
4	t
:	u
(v
,	w
n	x
m	y
8	z

Table 4.4: Random initial mapping used for Iteration 0

Listing 3: Code - Q5.(d)

```
1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import matplotlib.cm as cm
5 import matplotlib.colors as mcolors
6 from collections import Counter
7 from unicode import unicode
8 from operator import itemgetter
9 import random
10 import math
11 import csv
12
13 def text_cleaning(
14     symbol_file='symbols.txt',
15     message_file='message.txt',
16     training_file='war_and_peace_tolstoi.txt'
17 ):
18     # Loading the text, and estimating the transition probabilities
19
20     with open('war_and_peace_tolstoi.txt', 'r', encoding='utf-8') as file:
21         warpeace = file.read().lower()
22
23     with open('symbols.txt', 'r', encoding='utf-8') as symbol_file:
24         symbols_list = symbol_file.read().splitlines()
25
26     with open('message.txt', 'r', encoding='utf-8') as message_file:
27         message = message_file.read()
28
29     # Cleaning the text from all unwanted accents, but it makes new symbols appear
30     # so we also remove them.
31     warpeace_clean = unicode(warpeace)
32     warpeace_clean = warpeace_clean.replace('\n', ' ').replace('#', '')
33     warpeace_clean = warpeace_clean.replace('%', '').replace('$', '')
34
35     return symbols_list, message, warpeace_clean
36
37 def transition_counts(warpeace_clean):
38     # Keeping track of the number of counts of:
39     # 1) Each symbol
40     symbols_found = Counter()
41     # 2) Each pair of symbols
42     pair_found = Counter()
43
44     for i in range(len(warpeace_clean)-1):
45         s_i_minus_1 = warpeace_clean[i]
46         s_i = warpeace_clean[i+1]
47         pair_found[(s_i_minus_1, s_i)] += 1
```

```

48         symbols_found[s_i_minus_1] += 1
49
50     return symbols_found, pair_found
51
52 def psi_phi_calculations(symbols_found, pair_found, precision=1e-9):
53
54     transition_probabilities = {}
55
56     for (s_i_minus_1, s_i), frequency_pair in pair_found.items():
57         transition_probabilities[s_i_minus_1, s_i] = frequency_pair/(symbols_found[s_i_minus_1])
58
59     unsorted_symbol_states = set([i[0] for i in transition_probabilities.keys()] \
60 + [i[1] for i in transition_probabilities.keys()])
61     symbol_states = sorted(unsorted_symbol_states)
62     n_symbols = len(symbol_states)
63
64     symbol_indexes = {symbol: i for i, symbol in enumerate(symbol_states)}
65
66     # Transition probabilities matrix
67     psi = np.zeros((n_symbols, n_symbols))
68     for (s_i_minus_1, s_i), p_current_given_previous in transition_probabilities.items():
69         i = symbol_indexes[s_i_minus_1]
70         j = symbol_indexes[s_i]
71         psi[i, j] = p_current_given_previous
72
73     # Adapting for ergodicity, after answering question 5.(e)
74     for i in range(np.shape(psi)[0]):
75         for j in range(np.shape(psi)[0]):
76             if psi[i,j] == 0:
77                 psi[i,j] += precision
78
79
80     eigenvalues, eigenvectors = np.linalg.eig(psi.T)
81     stationary_vector = eigenvectors[:, np.isclose(eigenvalues, 1)]
82     stationary_distribution = (stationary_vector.real / np.sum(stationary_vector))
83
84     stationary_distribution_dict = \
85     {symbol: stationary_distribution[symbol_indexes[symbol]].real for symbol in symbol_states}
86
87     stationary_probabilities = [i for i in stationary_distribution_dict.values()]
88
89     return psi, stationary_distribution_dict, symbol_indexes, stationary_probabilities
90
91
92 def save_csv_data_transition(psi, filename="transition_probabilities.csv"):
93     psi_array = np.array(psi)
94     clean_data_transition = [[f"{value:.2e}" for value in row] for row in psi_array]
95
96     with open(filename, mode="w", newline="") as file:

```

```

97         writer = csv.writer(file)
98         writer.writerows(clean_data_transition)
99
100 def save_csv_data_stationary(
101     stationary_distribution_dict,
102     filename="stationary_probabilities.csv"
103 ):
104     with open(filename, mode="w", newline="") as file:
105         writer = csv.writer(file)
106         writer.writerow(["Symbol", "Stationary Probability"]) # Header row
107         for symbol, probability in stationary_distribution_dict.items():
108             clean_data=f"{probability:.3e}"
109             writer.writerow([symbol, clean_data])
110
111 # Summary tables - Q1.a)
112 def plot_heatmap_transition(psi, save=False):
113     plt.figure(figsize=(7, 7))
114     plt.imshow(psi, cmap="viridis", interpolation='nearest')
115     plt.colorbar(label='Probability')
116     plt.title('Pairs probabilities psi(alpha, beta)')
117     plt.xlabel("Symbol index")
118     plt.ylabel("Symbol index")
119     plt.tight_layout()
120     if save:
121         plt.savefig('heatmap_transition.png', format="png", dpi=300, bbox_inches="tight")
122     plt.show()
123
124
125 def plot_bar_stationary(symbols_list, stationary_probabilities, cmap_colors='winter', save=False):
126     stationary_probabilities = [float(prob) for prob in stationary_probabilities]
127     cmap = plt.get_cmap(cmap_colors, int(np.ceil(len(symbols_list))))
128     colors = [cmap(i) for i in np.linspace(0, 0.8, len(symbols_list))]
129
130     plt.figure(figsize=(14, 6))
131     plt.bar(symbols_list, stationary_probabilities, color=colors)
132     plt.title("Stationary distributions phi(gamma) of each symbol", fontsize=16)
133     plt.xlabel("Symbols", fontsize=14)
134     plt.ylabel("Stationary distribution probabilities", fontsize=14)
135     plt.tight_layout()
136     if save:
137         plt.savefig('bar_stationary.png', format="png", dpi=300, bbox_inches="tight")
138     plt.show()
139
140 # Q5.d)
141 def random_proposal(target_list):
142     indexes = np.arange(0, len(target_list)).tolist()
143     new_indexes = random.sample(indexes, len(indexes))
144     proposal_list = [target_list[i] for i in new_indexes]
145     proposal_mapping = {proposal_list[i] : target_list[i] for i in range(len(target_list))}

```

```
146     return proposal_mapping
147
148 def proposal_mechanism(mapping):
149     mapping_proposal = mapping.copy()
150     keys_list = list(mapping.keys())
151
152     indexes = np.arange(0, len(mapping)).tolist()
153     swap_index = random.sample(indexes, 2)
154
155     tmp = mapping_proposal[keys_list[swap_index[0]]]
156     mapping_proposal[keys_list[swap_index[0]]] = mapping_proposal[keys_list[swap_index[1]]]
157     mapping_proposal[keys_list[swap_index[1]]] = tmp
158
159     return mapping_proposal
160
161
162 def jointprob_e_given_sig(
163     message : str,
164     symbol_indexes : dict,
165     mapping : dict,
166     stationary_distribution_dict : list,
167     psi: np.ndarray
168 ):
169
170     log_jointprob_e = 0
171
172     reversed_symbol_indexes = {index: symbol for symbol, index in symbol_indexes.items()}
173
174     first_key = message[0]
175     inverse_mapping_first_key = mapping[first_key]
176     first_key_symbol = reversed_symbol_indexes[symbol_indexes[inverse_mapping_first_key]]
177
178     stationary_first_key = stationary_distribution_dict[first_key_symbol]
179
180     log_jointprob_e += math.log(stationary_first_key)
181
182     prev = message[0]
183     for i in range(1, len(message)):
184         prev = message[i-1]
185         current = message[i]
186         reverse_map_i_minus_1 = mapping[prev]
187         reverse_map_i = mapping[current]
188         index1, index2 = symbol_indexes[reverse_map_i_minus_1], symbol_indexes[reverse_map_i]
189
190         transition_prob_pair = psi[index1, index2]
191
192         if transition_prob_pair > 0:
193             log_jointprob_e += math.log(transition_prob_pair)
194         else:
```

```
195         log_jointprob_e += math.log(1e-10)
196
197     return log_jointprob_e
198
199 def mcmc_step(
200     current_mapping,
201     stationary_distribution_dict,
202     psi,
203     message,
204     symbol_indexes,
205 ):
206
207     accepted = False
208     new_mapping = proposal_mechanism(current_mapping)
209
210     log_joint_sig = jointprob_e_given_sig(message,
211                                           symbol_indexes,
212                                           current_mapping,
213                                           stationary_distribution_dict,
214                                           psi
215     )
216     log_joint_sig_prime = jointprob_e_given_sig(message,
217                                                  symbol_indexes,
218                                                  new_mapping,
219                                                  stationary_distribution_dict,
220                                                  psi
221     )
222
223     if log_joint_sig_prime > log_joint_sig:      # The ratio was not possible to compute without
224                                                  # introducing a runtime error.
225         A = 1
226     else:
227         A = math.exp(log_joint_sig_prime - log_joint_sig)
228
229     U_i = random.uniform(0, 1)
230
231     if U_i <= A:
232         returned_mapping = new_mapping
233         accepted = True
234     else:
235         returned_mapping = current_mapping
236         accepted = False
237
238     return returned_mapping, accepted
239
240
241 def decrypting_first60(mapping, message, end=60):
242
243     decrypted_message = ""
```

```
244     for char in message[:end]:
245         decrypted_char = mapping[char]
246         decrypted_message += decrypted_char
247
248     return decrypted_message
249
250
251 def run_mcmc(
252     n_iterations,
253     initial_mapping,
254     stationary_distribution_dict,
255     psi,
256     message,
257     symbol_indexes
258 ):
259
260     current_mapping = initial_mapping
261     mappings_list = [current_mapping]
262     n_accepted = 0
263
264     for i in range(n_iterations):
265         current_mapping, accepted = mcmc_step(current_mapping,
266                                             stationary_distribution_dict,
267                                             psi,
268                                             message,
269                                             symbol_indexes
270                                             )
271         mappings_list.append(current_mapping)
272         if i % 100 == 0:
273             print("Iteration ", i, " : ", decrypting_first60(current_mapping, message))
274         n_accepted += accepted
275
276     return mappings_list
277 def decrypting_message_final(message,
278                             symbols_list,
279                             stationary_distribution_dict,
280                             stationary_probabilities,
281                             psi,
282                             symbol_indexes,
283                             n_iterations=2000,
284                             random_map=False
285                             ):
286
287     proposal_encoding = random_proposal(symbols_list)
288
289     mapping_chain = run_mcmc(n_iterations,
290                             proposal_encoding,
291                             stationary_distribution_dict,
292                             psi,
```

```
293     message,
294     symbol_indexes)
295
296     return mapping_chain
297
298 def main():
299     # Opening all the files and cleaning the training text
300     symbols_list, message, warpeace_clean = text_cleaning()
301
302     # Counting the frequencies of each symbol and each pair
303     symbols_found, pair_found = transition_counts(warpeace_clean)
304     # Calculating the transition probability matrix, the stationary distribution's
305     # probabilities (stored in a dictionary), the symbol's indexes (in the original symbol file)
306     # and the corresponding stationary probabilities.
307     psi, stationary_distribution_dict, symbol_indexes, stationary_probabilities
308     = psi_phi_calculations(
309     symbols_found,
310     pair_found
311     )
312
313     # Explanatory plots to understand better the data
314     plot_heatmap_transition(psi)
315     plot_bar_stationary(symbols_list, stationary_probabilities)
316
317     # Saving the plots if wanted
318     save_plots = False
319     if save_plots:
320         save_csv_data_transition(psi, symbols_list)
321         save_csv_data_stationary(stationary_distribution_dict)
322
323     # Choosing the initial mapping as random
324     proposal_encoding = random_proposal(symbols_list)
325     # Decrypting the whole message with 30000 iterations
326     A = decrypting_message_final(message,
327     symbols_list,
328     stationary_distribution_dict,
329     stationary_probabilities,
330     psi,
331     symbol_indexes,
332     n_iterations=30000,
333     random_map=False)
334
335     # Decrypting the entire message and comparing it to the original.
336     print("Original encoding : ", decrypting_first60(mapping=proposal_encoding,
337     message=message,
338     end=(len(message)-1)))
339     print("Decoded text : ", decrypting_first60(mapping=A[-1],
340     message=message,
341     end=(len(message)-1)))
```



```
342  
343 # Running the main() function to observe all results  
344 main()
```

4.5 Question 5.(e)

By definition, if a Markov Chain is ergodic, it can reach a unique steady-state, independent of the starting point. To be ergodic, a Markov Chain must be:

- 1) aperiodic, *i.e.* $P(X_n = s | X_{n-1} = s) = p_{s \rightarrow s} > 0$, implying it must also have a non-zero diagonal. In other words, it must be able to stay in every state with a non-null probability.
- 2) irreducible, *i.e.* if there is a path (with non-zero probability) from each state to every other state in the transition graph.

In our case, if there are some zero entries for $\phi(\alpha, \beta)$, then our distribution is not irreducible nor aperiodic, meaning $p_{s \rightarrow s}$ is not > 0 . This makes sense since, in English, and especially in books such as War and Peace by Leo Tolstoy, there is a very low chance (if not a probability 0) to see a sequence of identical symbols (e.g. 'aaaaaaa', 'bbbb', ...) for every symbol (53 in total). Moreover, these sequences of identical symbols must be seen much more than once for their probabilities to be non-negligible compared to other transitions and thus not to be neglected by the computer's precision, e.g. in the book's summary, the chapter titles XII, XIII, etc. introduce a transition probability between some letters, but it might not be sufficient enough since it happens a few times only in the whole book. Regarding the irreducibility of the chain, the same reason applies: states cannot be reached from all possible other states ('xyz' never happens a lot in a well-written book).

A possible solution to restore the ergodicity of the chain is to make the transition probabilities non-zero. For instance, we could set a threshold (e.g. ranging from $1e-8$ to $1e-10$), and correct all transition probabilities (including the diagonal elements) below that threshold to make them equal to the threshold. It must be done before calculating the stationary distributions. This will force back the ergodicity of the chain by permitting a restoration of both aperiodicity and irreducibility of the chain. However, this threshold must be low enough and have no incidence on the other symbols probability of occurring, because the only goal of the threshold is to restore ergodicity and not replace symbols we are sure of.

4.6 Question 5.(f)

As stated for the previous question, this method is not perfect for decoding, and there are flaws in the approach we will study.

The approach relies on analyzing how english words are formed to obtain a sequence of symbols making sense. Symbol probabilities alone won't be sufficient, we might obtain a sequence of spaces with a few of most-used letters in the english language (a, e, etc.), and there might even be no letters at all since the space probability is very high compared to them. The sequences of symbols won't make sense and the program won't fulfill the initial task, a sufficient reason to refute this idea.

Using a second-order Markov chain can introduce some computation issues, because the program will have to work in three dimensions for transition probabilities (one for each symbol s_i, s_{i-1} and s_{i-2}). This makes the computation much more complex because the transition probability matrix will have $N \times N \times N$ in addition to a third dimension. Since eigenvalues are calculated for 2D matrices, we might also have to adapt the code to fit 3D eigenvalues/eigenvectors methods. An additional information to refute this idea of second-order Markov chain is the fact the training data must be much higher for the model to be as accurate as a first-order Markov Chain.

If the encryption scheme allows two symbols to be mapped to the same encrypted value, the approach will simply first attribute spaces, and then decrypt the most used word in the training set that has the same number of letters. We will obtain a sentence which won't make sense, and thus it will not decrypt the message as intended. Or we could even simply have a sequence of random spaces and letters, since

only depending on stationary distributions.

If we use the approach on Chinese language for instance, it is highly probable the program won't be able to decrypt the message. Firstly, this is because Chinese signs are unique and don't work as English and other Western languages where letters and symbols represent how we pronounce the word. Chinese language associates images (each having a full meaning alone) to each symbol, which makes their use much less frequent in sentences than letters in English and will be much more complicated for the computer to grasp. For example, if the sentence to decrypt is "The dog eats outside.", the computer won't make the difference between "cat" and "dog", between "eating" and "playing", etc. It will likely decrypt the wrong sentence. Additionally, their transition probabilities and stationary probabilities might end up to be very low values, which will be harder for the program to decrypt. Using a two-symbol swap as we did in the approach is also too low, and the time for the program to decrypt the language will be too higher. Hence, to resume what will happen knowing all this information: the program will offer a two symbol swap; the difference in probability will be so low, that they will most likely be rejected during the acceptance step of the Metropolis-Hasting algorithm. We will end-up in an infinite loop.

A remark we can make when displaying the whole decrypted message is the fact some letters still can't be decrypted right by the program since they are not very used (e.g. j, q, z, k and probably x), even after a high number of iterations. What we can see in common for these letters is their very low stationary probabilities. Their probabilities are somehow lower or equal to punctuation probabilities. We could also apply the threshold method to the stationary vector, but I doubt it will solve the issue, because they might have an impact on highly-probable symbols by replacing them where they were expected.

5 | Question 7 - Optimization

5.1 Question 7.(a)

We want to find the local extrema of the function $f(x, y) = x + 2y$ subject to the constraints $y^2 + xy = 1$. This is a constrained optimization problem, with an equality constraint.

We name $g(x, y) = y^2 + xy - 1$, where finding the values of x and y such that $g(x, y) = 0$ is a reformulation of the constraint of our problem. Introducing a Lagrange multiplier, denoted by λ , the problem is equivalent to solving the following system of equations:

$$\begin{cases} \nabla f(x, y) + \lambda \nabla g(x, y) = 0 & (E_1) \\ g(x, y) = 0 & (E_2) \end{cases} \quad (5.1)$$

∇f and ∇g are in three dimensions, one for each of the following variables: x, y, λ . Differentiating with respect to each variables the equations, we get:

$$\frac{\partial E_1}{\partial x} = 0 \iff \frac{\partial}{\partial x} (x + 2y + \lambda(y^2 + xy - 1)) = 0 \iff \lambda y + 1 = 0 \quad (5.2)$$

Calculated a similar manner:

$$\frac{\partial E_1}{\partial y} = 0 \iff 2 + 2\lambda y + \lambda x = 0 \quad (5.3)$$

$$\frac{\partial E_1}{\partial \lambda} = 0 \iff y^2 + xy - 1 = 0 \quad (5.4)$$

Thus:

$$\begin{cases} \lambda y + 1 = 0 \\ 2 + 2\lambda y + \lambda x = 0 \\ y^2 + xy - 1 = 0 \end{cases} \quad (5.5)$$

Equivalent to:

$$\begin{cases} \lambda = -\frac{1}{y} \\ 2 + 2\left(-\frac{1}{y}\right)y + \left(-\frac{1}{y}\right)x = 0 \\ y^2 + xy - 1 = 0 \end{cases} \quad \begin{cases} \lambda = -\frac{1}{y} \\ -\frac{x}{y} = 0 \\ y^2 + xy - 1 = 0 \end{cases} \quad \begin{cases} \lambda = -\frac{1}{y} \\ x = 0 \\ y^2 - 1 = 0 \end{cases}$$

We obtain as results $y = \pm 1$, $x = 0$, and $\lambda = \mp 1$. Thus, we obtained the solutions to the system and to the optimization problem, *i.e.* the locations of the local extrema:

$$(x, y) \in \mathcal{S} = \{(0, 1), (0, -1)\}$$

5.2 Question 7.(b)

We assume we dispose of a method to evaluate the exponential function $\exp(x) = e^x$. We would like to evaluate the function $\ln a$, for a given $a \in \mathbb{R}_+$ using Newton's method.

5.2.1 Question 7.(b).(i)

Firstly, we can recall the exponential function's inverse mapping is the logarithmic function, *i.e.* $\exp(\ln x) = x$, and we can additionally write $\exp^{-1}x = \ln x$.

Thus, we can use the following function $f(x, a)$ to which Newton's Method can be applied to x such that $x = \ln a$:

$$f(x, a) = e^x - a \quad (5.6)$$

And by solving the equation $f(x, a) = 0$, we obtain the logarithm of a , *i.e.* $\ln a$.

5.2.2 Question 7.(b).(ii)

The update equation in Newton's method to search for the root of f , $f(x, a) = 0$, is found using the following formula:

$$x_{n+1} = x_n - \frac{f(x_n, a)}{f'(x_n, a)} \quad (5.7)$$

Calculating the first derivative of f with respect to x_n :

$$\begin{aligned} f'(x_n, a) &= \frac{\partial (e^{x_n} - a)}{\partial x_n} \\ f'(x_n, a) &= e^{x_n} \end{aligned} \quad (5.8)$$

Therefore, we obtain:

$$x_{n+1} = x_n - \frac{e^{x_n} - a}{e^{x_n}} \quad \Longleftrightarrow \quad x_{n+1} = x_n - 1 + \frac{1}{e^{x_n}} \quad (5.9)$$

6 | Question 8 - [BONUS] Eigenvalues as solutions of an optimization problem

We define A as a symmetric $n \times n$ - matrix, and:

$$q_A(x) = x^T A x \quad \text{and} \quad R_A(x) = \frac{q_A(x)}{\|x\|^2} = \frac{x^T A x}{x^T x} \quad \text{for } x \in \mathbb{R}^n. \quad (6.1)$$

We remind the purpose of this problem is to verify the fact: *If A is a symmetric $n \times n$ -matrix, the optimization problem $x^* = \arg \max_{x \in \mathbb{R}^n} R_A(x)$ has a solution, $R_A(x^*)$ is the largest eigenvalue of A , and x^* is a corresponding eigenvector.*

6.1 Question 8.(a)

The **Extreme Value theorem** states the following:

If a function f is continuous on a closed interval $[a; b]$, then f attains both a maximum and a minimum in that interval.

$$i.e. \quad \exists x_{max} \in [a; b] \quad / \quad \forall x \in [a; b] \quad f_{max} \geq f(x) \quad (6.2)$$

$$\text{and} \\ \exists x_{min} \in [a; b] \quad / \quad \forall x \in [a; b] \quad f_{min} \leq f(x) \quad (6.3)$$

The conditions for the validity of such theorem are requirements into the function f that f should be both **continuous** and **defined on a closed interval** (the function must be "compact").

We would like to show that $\sup_{x \in \mathbb{R}^n} R_A(x)$ using the extreme value theorem. Firstly, since \mathbb{R}^n is not compact, we need to find an interval of definition to be compact while having an equivalent supremum. Let's prove the unit sphere $S = \{x \in \mathbb{R}^n \mid \|x\| = 1\}$, which is a compact set, has an equivalent supremum.

Then, defining $s \in \mathbb{R}^n$, we notice $\forall s \in \mathbb{R}^n, \exists x \in S \quad / \quad x = \frac{s}{\|s\|}$. This expression is true, because by dividing by the norm of each vector in the space \mathbb{R}^n , we obtain a unit vector comprised in the interval S . By replacing in $R_A(x)$'s expression, we get:

$$\begin{aligned} \sup_{\{x \in S\}} R_A(x) &= \sup_{\{x \in \mathbb{R}^n \mid \|x\|=1\}} \frac{x^T A x}{x^T x} \\ &= \sup_{\left\{ \frac{s}{\|s\|} \in \mathbb{R}^n \mid \left\| \frac{s}{\|s\|} \right\| = 1 \right\}} \frac{\left(\frac{s}{\|s\|} \right)^T A \frac{s}{\|s\|}}{\left(\frac{s}{\|s\|} \right)^T \frac{s}{\|s\|}} \\ &= \sup_{\left\{ \frac{s}{\|s\|} \in \mathbb{R}^n \mid \left\| \frac{s}{\|s\|} \right\| = 1 \right\}} \frac{\frac{s^T}{\|s\|} A \frac{s}{\|s\|}}{\frac{s^T}{\|s\|} \frac{s}{\|s\|}} \end{aligned}$$

And since $\|s\|$ is a scalar:

$$\begin{aligned}
 &= \sup_{\left\{ \frac{s}{\|s\|} \in \mathbb{R}^n \mid \left\| \frac{s}{\|s\|} \right\| = 1 \right\}} \frac{\frac{s^T}{\|s\|^2} A s^T}{\frac{s^T s}{\|s\|^2}} \\
 &= \sup_{\left\{ \frac{s}{\|s\|} \in \mathbb{R}^n \mid \left\| \frac{s}{\|s\|} \right\| = 1 \right\}} \frac{s^T A s}{s^T s}
 \end{aligned}$$

And switching back to the domain of definition of s , since the :

$$\begin{aligned}
 &= \sup_{\{s \in \mathbb{R}^n\}} \frac{\frac{s^T}{\|s\|^2} A s^T}{\frac{s^T s}{\|s\|^2}} \\
 &= \sup_{\{s \in \mathbb{R}^n\}} R_A(s)
 \end{aligned}$$

Leading to the final expression of:

$$\sup_{x \in S} R_A(x) = \sup_{s \in \mathbb{R}^n} R_A(s) \quad (6.4)$$

Which proves the supremum of the set $S = \{x \in \mathbb{R}^n \mid \|x\| = 1\}$ is equivalent to that of \mathbb{R}^n .

Since the set S is continuous (because its expression, the unit sphere, is valid for all element of \mathbb{R}^n), and is bounded (its images limits are finite numbers, *i.e.* all equal to 1), then the **Extreme Value theorem applies** to our case and therefore, the supremum of $R_A(x), x \in \mathbb{R}^n$ is attained when the supremum of $R_A(x_s), x_s \in S$ is attained.

6.2 Question 8.(b)

We define $\lambda_1 \geq \dots \geq \lambda_n$ as the eigenvalues of A in descending order, and ξ_1, \dots, ξ_n their corresponding eigenvectors that form an ONB.

Using our previous results, we start from the unit sphere:

$$\sup_{x \in \mathbb{R}^n \mid \|x\|=1} \frac{x^T A x}{\|x\|} = \sup_{x \in \mathbb{R}^n \mid \|x\|=1} x^T A x \quad (6.5)$$

For all $x \in \mathbb{R}^n$, we can rewrite using A 's eigenvectors (since they form an ONB) as:

$$x = \sum_{i=1}^n (\xi_i^T x) \xi_i \quad (6.6)$$

Thus, replacing in our problem, we get:

$$x^T A x = \left(\sum_{i=1}^n (\xi_i^T x) \xi_i \right)^T A \left(\sum_{i=1}^n (\xi_i^T x) \xi_i \right) = \left(\sum_{i=1}^n ((\xi_i^T x) \xi_i)^T \right) A \sum_{i=1}^n (\xi_i^T x) \xi_i$$

If we recall the following property of eigenvectors of the matrix A :

$$A\xi_i = \lambda_i \xi_i$$

Then since A is symmetric, the property is valid for the transpose of eigenvectors:

$$A\xi_i^T = \lambda_i \xi_i^T$$

Therefore:

$$x^T A x = \left(\sum_{i=1}^n ((\xi_i^T x) \xi_i)^T \right) \sum_{i=1}^n A \xi_i^T x \xi_i = \left(\sum_{i=1}^n ((\xi_i^T x) \xi_i)^T \right) \sum_{i=1}^n \lambda_i \xi_i^T x \xi_i$$

Using the property of transpose:

$$(AB)^T = B^T A^T$$

Then:

$$\begin{aligned} x^T A x &= \left(\sum_{i=1}^n \xi_i^T (\xi_i^T x)^T \right) \sum_{i=1}^n \lambda_i \xi_i^T x \xi_i = \left(\sum_{i=1}^n \xi_i^T x^T (\xi_i^T)^T \right) \sum_{i=1}^n \lambda_i \xi_i^T x \xi_i \\ &= \left(\sum_{i=1}^n \xi_i^T x^T \xi_i \right) \sum_{i=1}^n \lambda_i \xi_i^T x \xi_i \\ x^T A x &= \left(\sum_{i=1}^n \xi_i^T x^T \xi_i \right) \left(\lambda_i \sum_{i=1}^n \xi_i^T x \xi_i \right) \end{aligned} \quad (6.7)$$

Since λ_i is a scalar.

And, since $\xi_i, i \in \{1, \dots, n\}$ form an orthonormal basis (ONB), then:

$$\xi_i^T \xi_j = \begin{cases} 1, & i = j \\ 0, & i \neq j \end{cases} \quad (6.8)$$

Thus:

$$\begin{aligned} x^T A x &= \sum_{i=1}^n \lambda_i \xi_i^T x^T x \xi_i \\ x^T A x &= \sum_{i=1}^n \lambda_i x_i^T \xi_i \end{aligned}$$

And once again using the property of the ONB vectors 6.8, then:

$$x^T A x = \sum_{i=1}^n \lambda_i \quad (6.9)$$

Since we defined the eigenvalues of A previously as $\lambda_1 \geq \dots \geq \lambda_n$ in descending order, then:

$$R_A(x \in S) = \sum_{i=1}^n \lambda_i \leq \lambda_1 \quad (6.10)$$

Which was what we intended to prove.

6.3 Question 8.(c)

We consider $x \in \mathbb{R}^n \setminus \text{span}\{\xi_1, \dots, \xi_k\}$, which means all x in \mathbb{R}^n that is not a part of $\text{span}\{\xi_1, \dots, \xi_k\}$. The $\text{span}\{\xi_1, \dots, \xi_k\}$ contains all the linear combinations of linearly independent eigenvectors corresponding to λ_1 , $k \leq n$.

Therefore, x can be rewritten as we did before for Question 8.(b), but excluding all the eigenvectors corresponding to λ_1 :

$$x = \sum_{i=k+1}^n (\xi_i^T x) \xi_i \quad (6.11)$$

Using our previous result for Question 8.(b), we can state (remembering the eigenvalues are ordered in descending order and thus λ_{k+1} is the maximum):

$$R_A(x) = x^T A x = \sum_{i=k+1}^n \lambda_i \leq \lambda_{k+1}$$

And since $\lambda_{k+1} < \lambda_1$, then:

$$R_A(x) < \lambda_1 x \in \mathbb{R}^n \setminus \text{span}\{\xi_1, \dots, \xi_k\} \quad (6.12)$$

Which is what we needed to prove.